

Artistic Stylization of Images and Video

Part III – Anisotropy and Filtering

Eurographics 2011

Jan Eric Kyprianidis

Hasso-Plattner-Institut, University of Potsdam, Germany



- **Stylized Augmented Reality for Improved Immersion**
Fischer et al., 2005
- **Real-time Video Abstraction**
Winnemöller et al., SIGGRAPH 2006
- **Coherent Line Drawings**
Kang et al., NPAR 2007
- **Structure Adaptive Image Abstraction**
Kyprianidis & Döllner, EG Theory and Practice of Computer Graphics 2008
- **Flow-based Image Abstraction**
Kang et al., Transactions on Visualization and Computer Graphics 2009
- **Artistic Edge and Corner Preserving Smoothing**
Papari et al., IEEE Transactions on Image Processing 2007
- **Image and Video Abstraction by Anisotropic Kuwahara Filtering**
Kyprianidis et al., Pacific Graphics 2009
- **Shape-simplifying Image Abstraction**
Kang & Lee, Pacific Graphics 2008
- **Image and Video Abstraction by Coherence-Enhancing Filtering**
Kyprianidis & Kang, Eurographics 2011

Non-photorealistic display of both the camera image and virtual objects:

- **Abstraction:** Bilateral filter applied to Gaussian pyramid and then upsampled
- **Edges:** Canny edge detector + morphological dilation



Conventional augmented reality

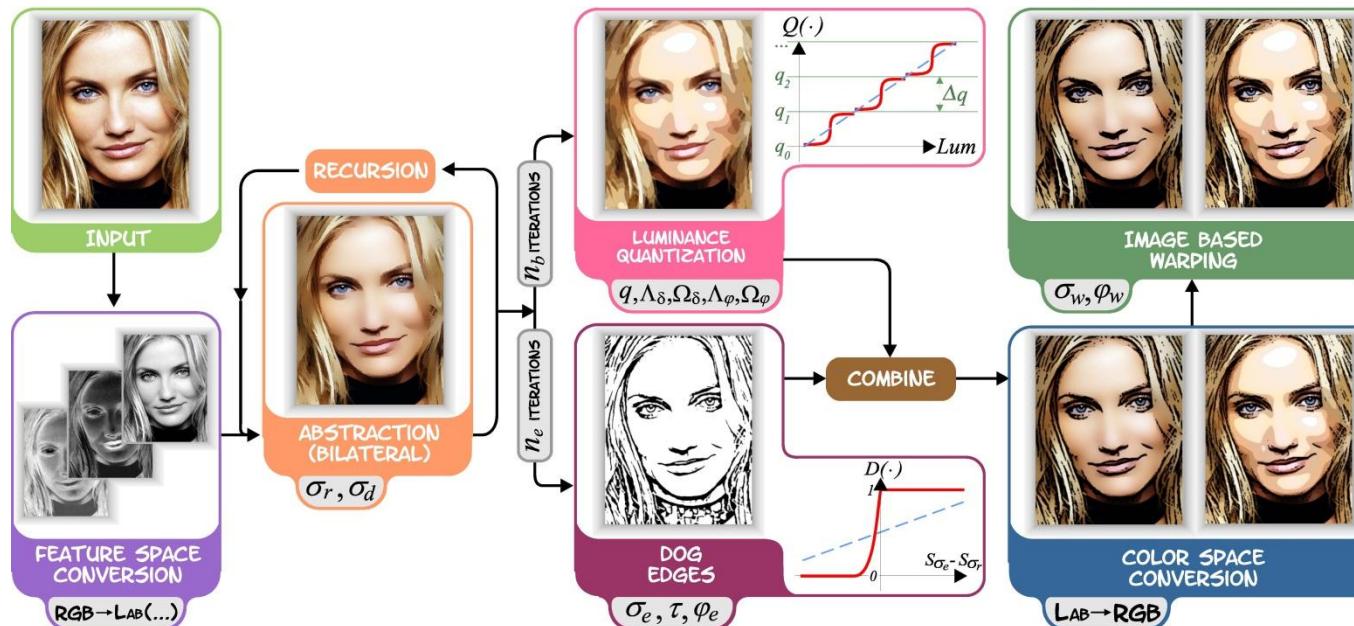


Stylized augmented reality

Image credit: Fischer et. al. (2005)

- **Abstraction:** Multiple iterations of xy-separable bilateral filter + color quantization
- **Edges:** Difference of Gaussians + thresholding

Image credit: Winnemöller et. al. (2006)





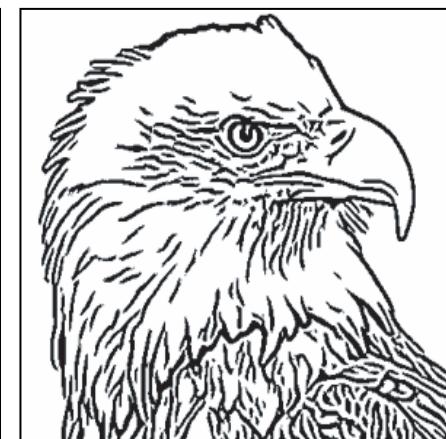
- **Edges:** 1D difference of Gaussians directed by flow field + flow-guided smoothing and thresholding.



Input image



Edge Tangent Flow



Line drawing



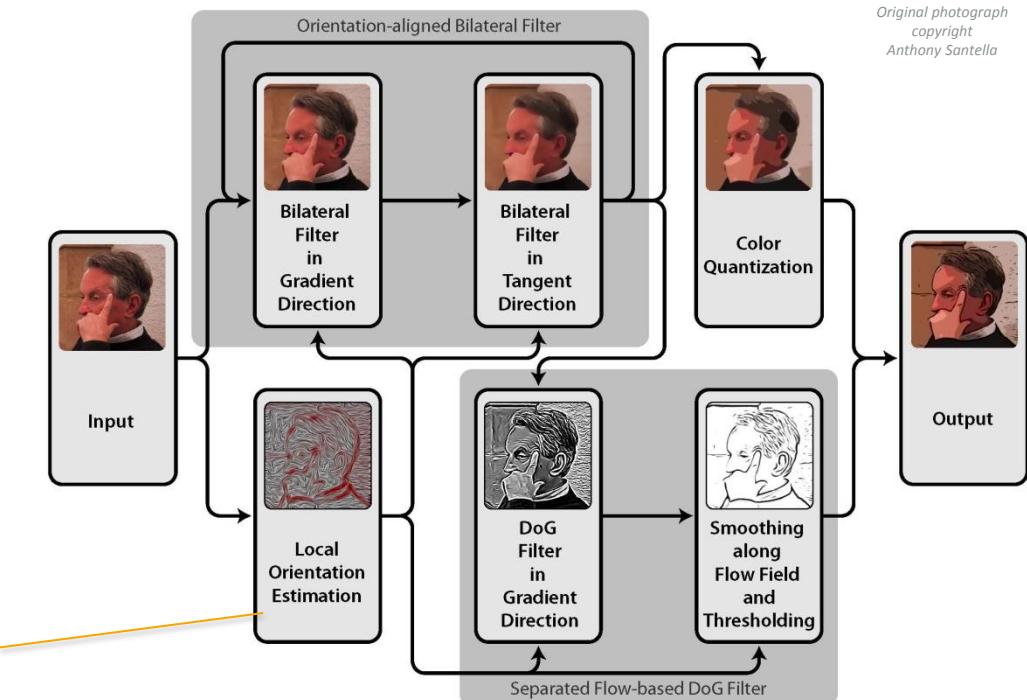
Flow-based filtering

Structure Adaptive Image Abstraction

Kyprianidis & Döllner (2008)

- **Abstraction:** Multiple iterations of orientation-aligned bilateral filter
- **Edges:** separable flow-based difference of Gaussians

Local orientation and an anisotropy measure derived from the smoothed structure tensor are used to guide the bilateral and difference of Gaussians filters

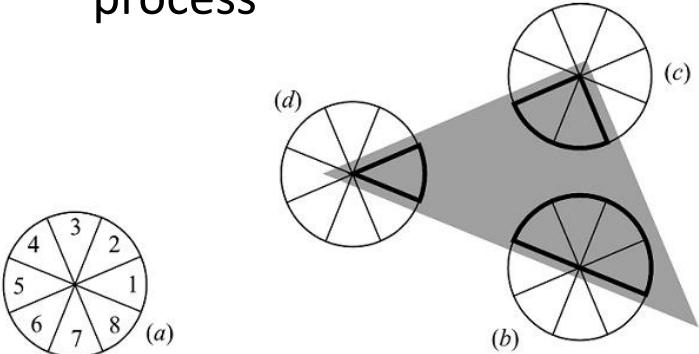


- **Abstraction:** Multiple iterations of flow-based bilateral filter
- **Edges:** (separable) flow-based difference of Gaussians
- Local orientation estimation of both techniques is based on the edge tangent flow (ETF)



Image credit: Kang et. al. (2009)

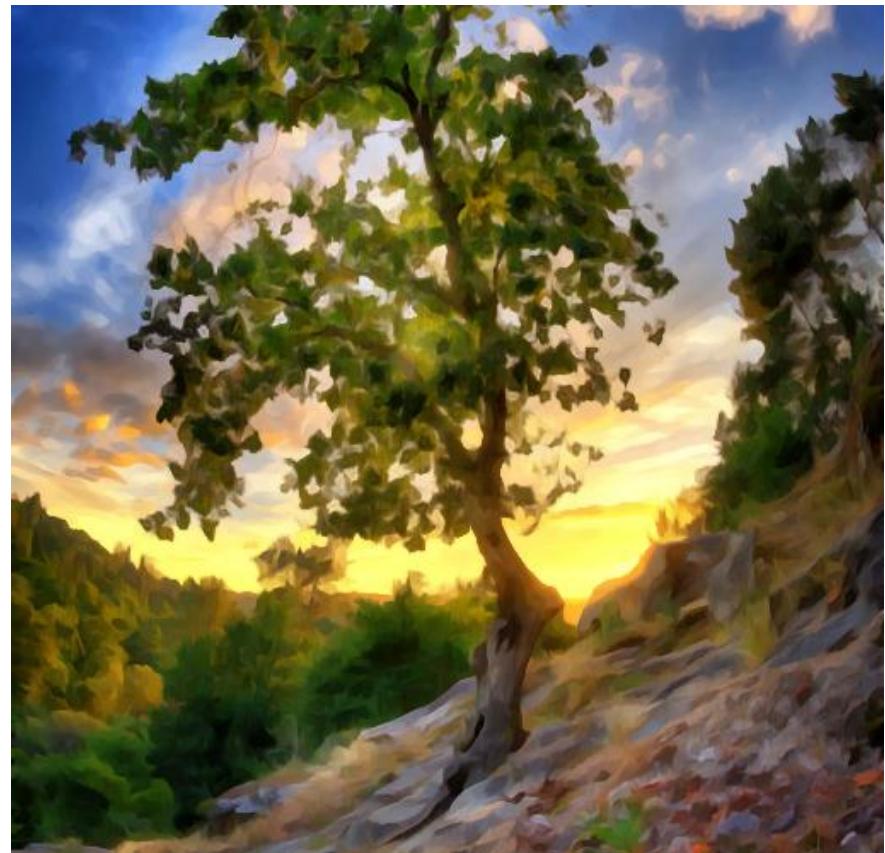
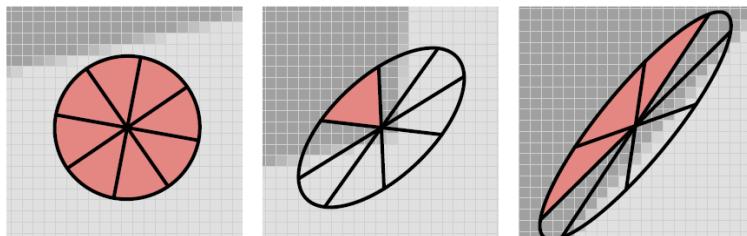
- Generalization of the Kuwahara filter. Creates output with a painterly look.
- Addresses two key issues of the original Kuwahara filter:
 - Rectangular subregions
 - Unstable subregion selection process



Credit for images: Papari et. al. (2009)



- Further generalization of the Kuwahara filter.
- Adaptation of the shape, scale and orientation of the filter to the local image structure.



Original image by Paulo Brandão@flickr.com



- PDE-based technique that simultaneous simplifies colors and shape:
 - Constrained mean curvature flow
 - Shock filter

Image credit: Kang & Lee (2008) / original image by Tambako the Jaguar@flickr.com



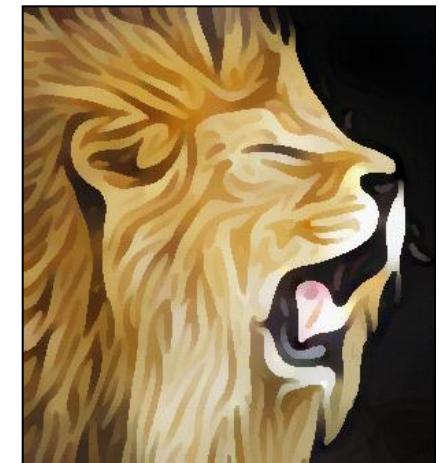
Input



20 iterations



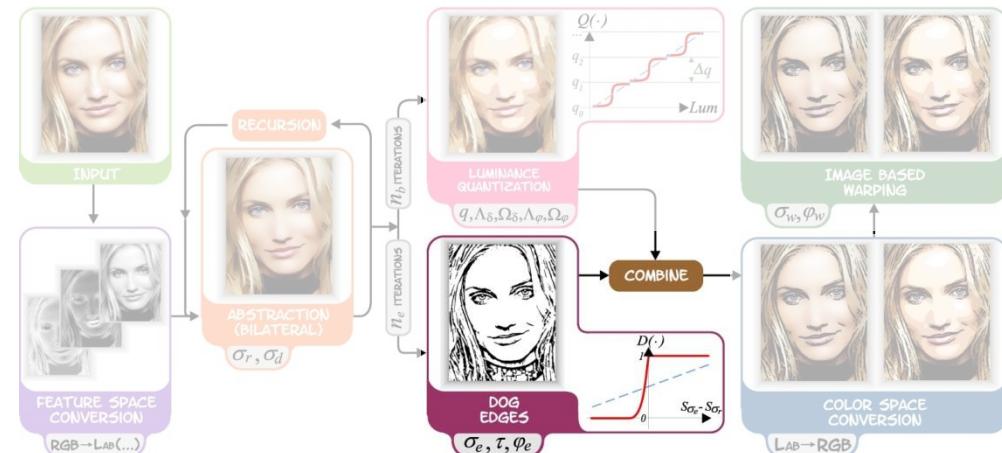
40 iterations



60 iterations

Difference of Gaussians:

- Laplacian of Gaussian (LoG)
- Isotropic Difference of Gaussians (DoG)
- Flow-based Difference of Gaussians
- Separable Flow-based Difference of Gaussians





DoG Edges vs Canny Edges

Original image from USC-SIPI Image Database



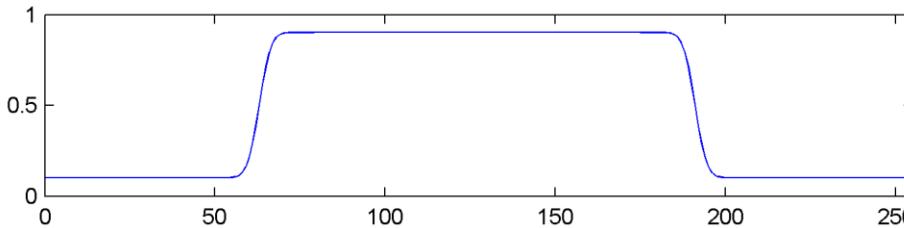
Canny Edges



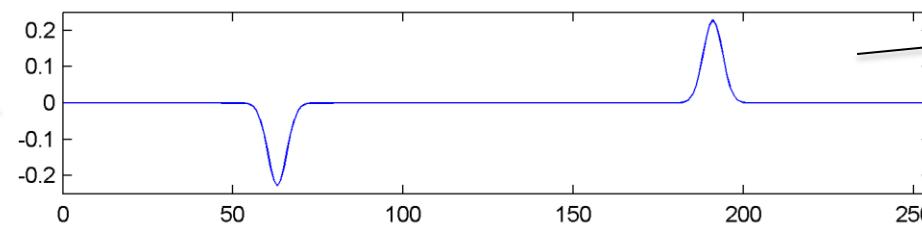
Flow-based difference of Gaussians

Edge profile without noise:

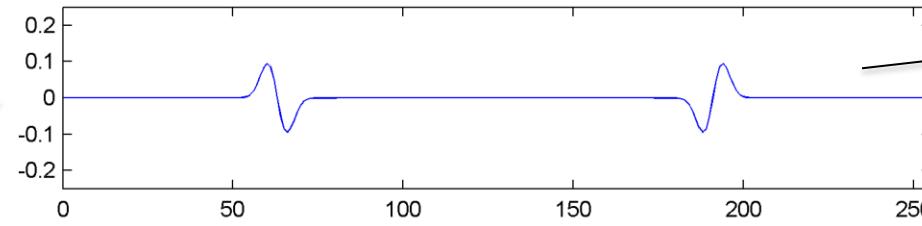
Scanline of an image



First derivative of scanline



Second derivative of scanline

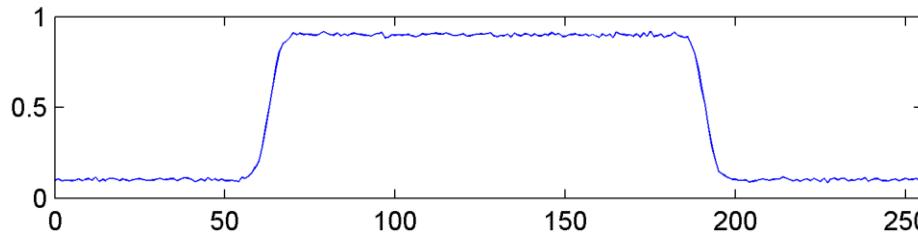


Edges are located at minima or maxima of the first derivative

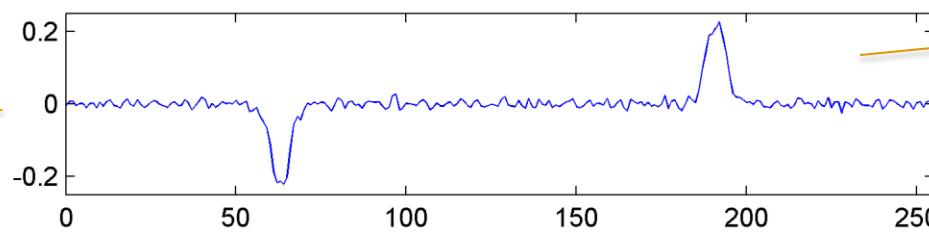
Edges are located at zero crossings of the second derivative

Edge profile with noise:

Scanline of an image with noise

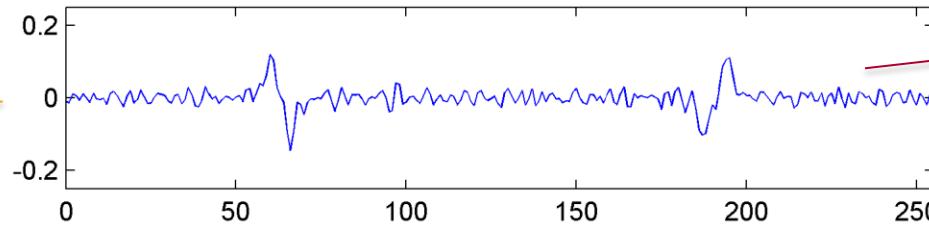


First derivative of scanline



First derivative sensitive to noise

Second derivative of scanline



Second derivative even more sensitive to noise

In 2D the second derivative corresponds to the Laplacian:

$$L = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

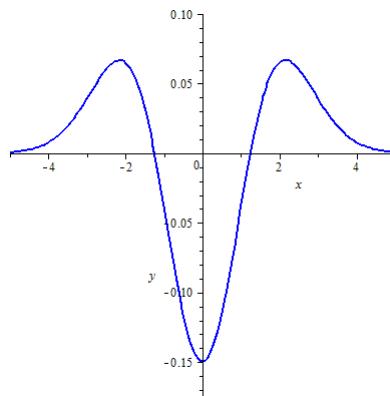
Similar to the second derivative the Laplacian is sensitive to noise. To make the Laplacian less sensitive to noise, apply a Gaussian to the image first:

$$\text{LoG} = L \star G_\sigma ,$$

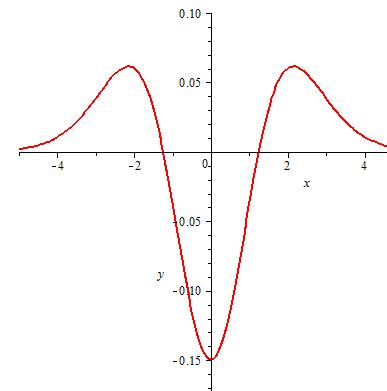
where G_σ is a 2D Gaussian with standard deviation σ :

$$G_\sigma(x, y) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

A Laplacian of Gaussian can be approximated by a difference of Gaussians:



$$\text{LoG}(x, y) = L \star G_\sigma$$



$$\text{DoG}(x, y) = G_{\sigma_i}(x, y) - G_{\sigma_e}(x, y)$$

Good engineering
solution:
 $\sigma_i = 1.6 \cdot \sigma_e$

Fast to implement
since Gaussian
is separable

Zero-crossing are found by thresholding:



An approach to create smooth edges was proposed by Winnemöller et al.:

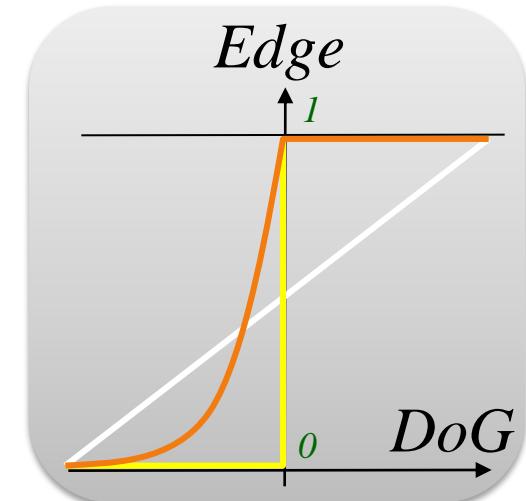
$$D(\sigma_e, \tau, \varphi_e) = \begin{cases} 1 & \text{if } (G_{\sigma_e} - \tau G_{1.6 \cdot \sigma_e}) > 0 \\ 1 + \tanh(\varphi_e \cdot G_{\sigma_e} - \tau G_{1.6 \cdot \sigma_e}) & \text{otherwise} \end{cases}$$

- The parameter τ controls the sensitivity to noise. A typical values are $\tau = 0.98$ or $\tau = 0.99$.
- The falloff parameter φ_e determines the sharpness of edge representations, typical values are $\varphi_e \in [0.75, 5.0]$.



Credit for slide: H. Winnemöller

\tanh
(half-truncated)

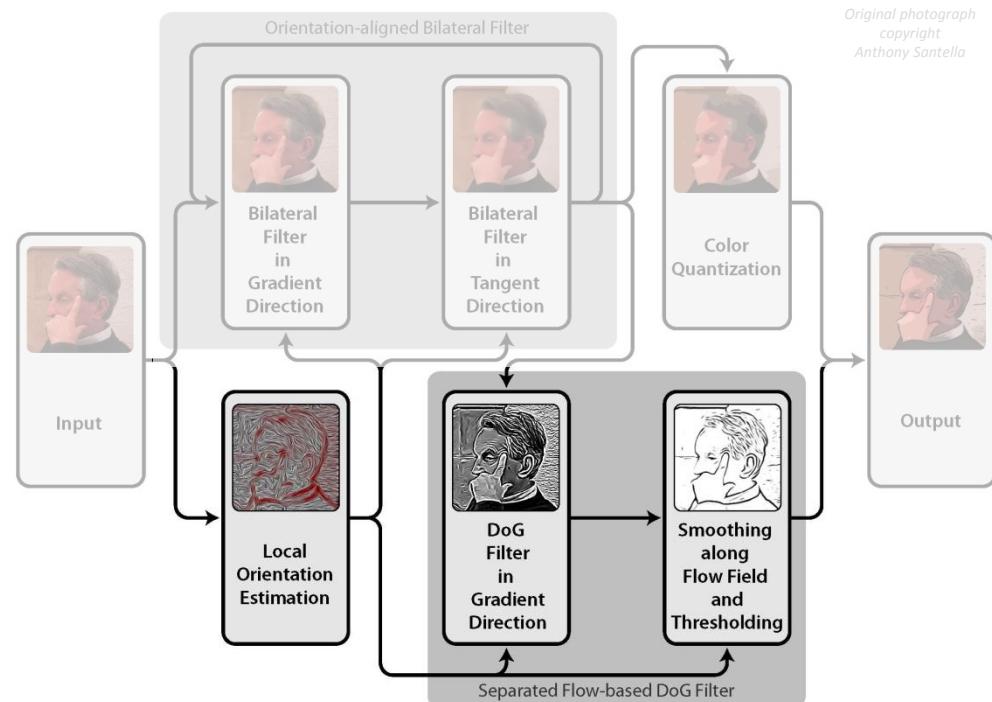


Local Structure Estimation:

- Edge Tangent Flow
- Structure Tensor

DoG Guided by Local Image Structure:

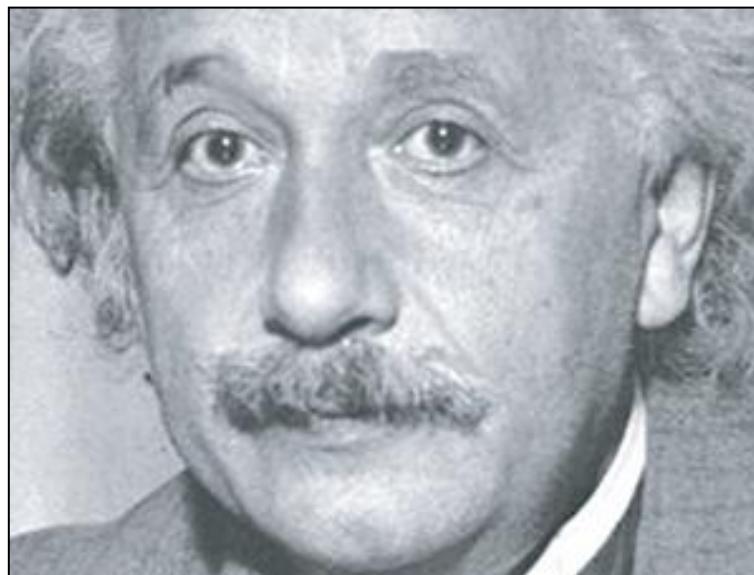
- Flow-based DoG
- Separable flow-based DoG





Edge Tangent Flow (ETF):

- Smoothly varying vector field
- Feature-preserving flow



Input image



Edge Tangent Flow

Image credit: Kang et al. (2007)



Weighted vector smoothing similar to bilateral filter:

Multiple iterations (≈ 3)

$$t^{n+1}(x) = \frac{1}{k} \sum_{y \in \Omega(x)} \phi(x, y) \cdot t^n(y) \cdot w_s(x, y) \cdot w_m(x, y) \cdot w_d(x, y)$$

t^0 is calculated using the Sobel filter.

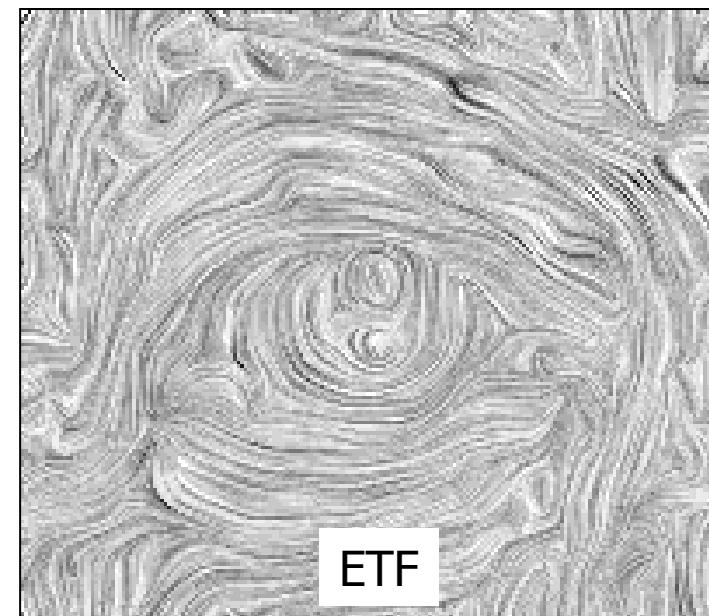
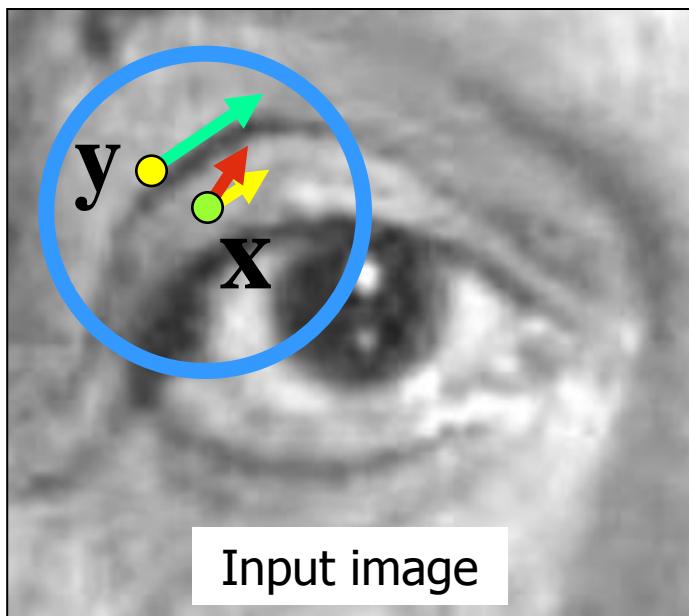


Image credit: Kang et al. (2007)



$$t^{n+1}(x) = \frac{1}{k} \sum_{y \in \Omega(x)} \phi(x, y) \cdot t^n(y) \cdot w_s(x, y) \cdot w_m(x, y) \cdot w_d(x, y)$$

Assure different vectors point in the same direction

$$\phi(x, y) = \text{sign}(t^n(x) \cdot t^n(y))$$

$$w_s(x, y) = \begin{cases} 1 & |x - y| < r \\ 0 & \text{else} \end{cases}$$

Restrict filtering to a predefined radius

More weight to vectors with higher gradient magnitude

$$w_m(x, y) = \frac{1}{2} [1 + \tanh(|g(x)| - |g(y)|)]$$

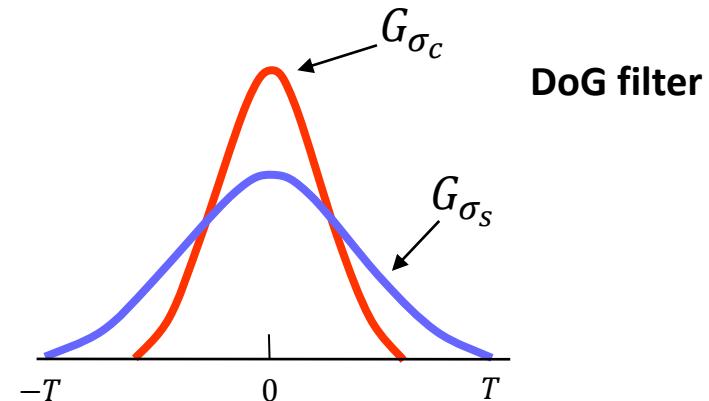
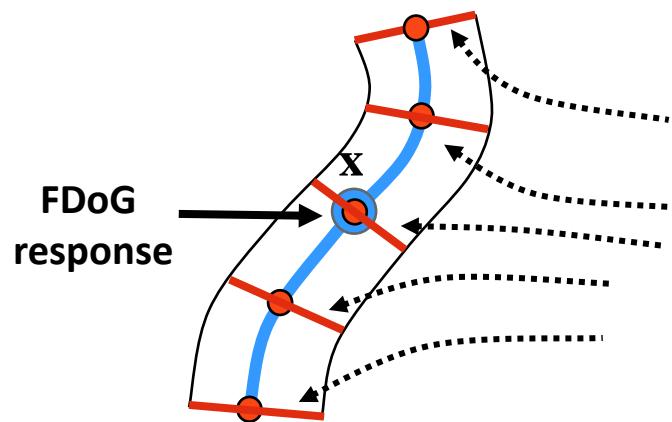
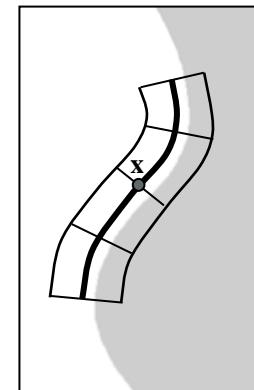
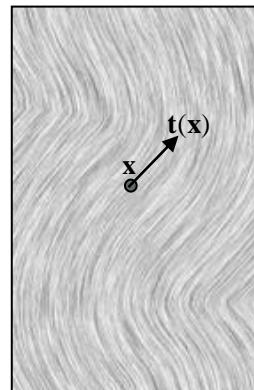
$$w_d(x, y) = |t^n(x) \cdot t^n(y)|$$

More weight for vectors with direction similar to current filter origin

Flow-based Difference of Gaussians

Kang et al. (2007)

Image credit: Kang et al. (2007)



Let $f: \mathbb{R}^2 \rightarrow \mathbb{R}^3$ denote the input image and let

$$\frac{\partial f}{\partial x} = \begin{pmatrix} \frac{\partial R}{\partial x} & \frac{\partial G}{\partial x} & \frac{\partial B}{\partial x} \end{pmatrix}^t \quad \frac{\partial f}{\partial y} = \begin{pmatrix} \frac{\partial R}{\partial y} & \frac{\partial G}{\partial y} & \frac{\partial B}{\partial y} \end{pmatrix}^t$$

be the partial derivatives of f .

The structure tensor is then defined by:

$$(g_{ij}) = J^t J = \begin{pmatrix} \left\langle \frac{\partial f}{\partial x}, \frac{\partial f}{\partial x} \right\rangle & \left\langle \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right\rangle \\ \left\langle \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right\rangle & \left\langle \frac{\partial f}{\partial y}, \frac{\partial f}{\partial y} \right\rangle \end{pmatrix} =: \begin{pmatrix} E & F \\ F & G \end{pmatrix}$$

The structure tensor is a 2×2 symmetric positive semidefinite matrix

These can be implemented for example using Gaussian derivatives or the Sobel filter.

In differential geometry the structure tensor is also known as first fundamental form

The induced quadratic form of the structure tensor measures the squared rate of change of f in direction of a vector $n = (n_x, n_y)$:

$$S(n) = En_x^2 + 2Fn_xn_y + Gn_y^2$$

The extremal values of $S(n)$ on the unit circle correspond to the eigenvalues of (g_{ij}) :

$$\lambda_{1,2} = \frac{E + G \pm \sqrt{(E - G)^2 + 4F^2}}{2}$$

The corresponding eigenvectors are:

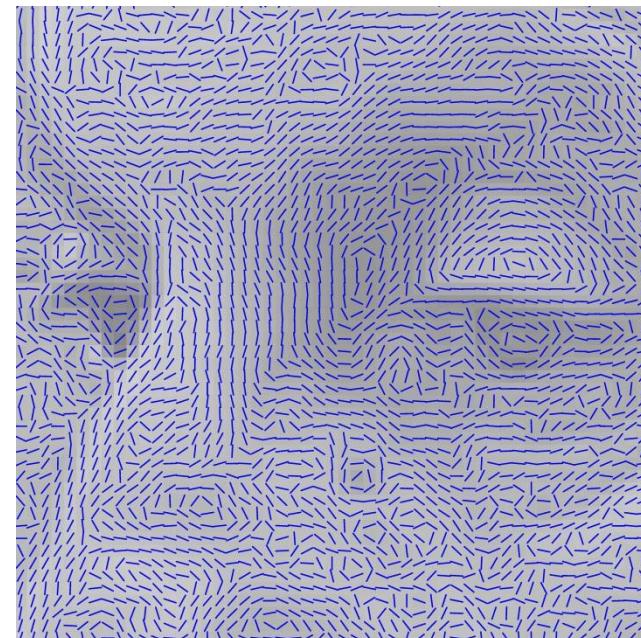
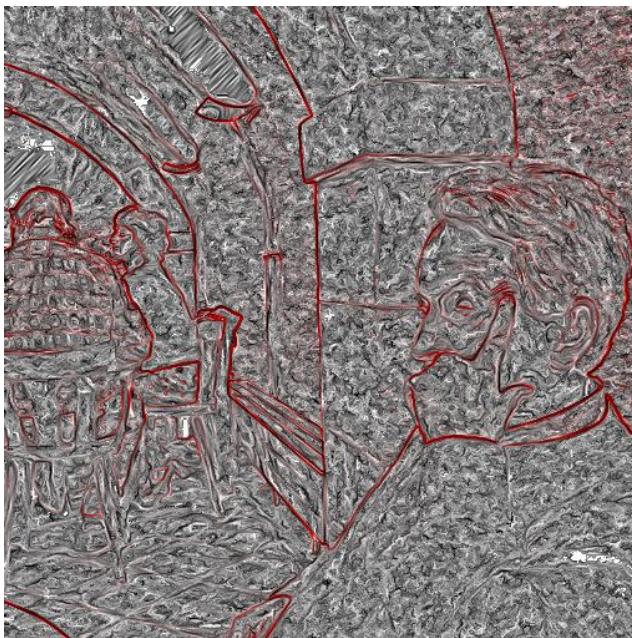
Eigenvector of major eigenvalue.
Direction of maximum change:
gradient direction.

$$v_1 = \begin{pmatrix} F \\ \lambda_1 - E \end{pmatrix} \quad v_2 = \begin{pmatrix} \lambda_1 - E \\ -F \end{pmatrix}$$

Eigenvector of minor eigenvalue.
Direction of minimum change:
tangent direction.

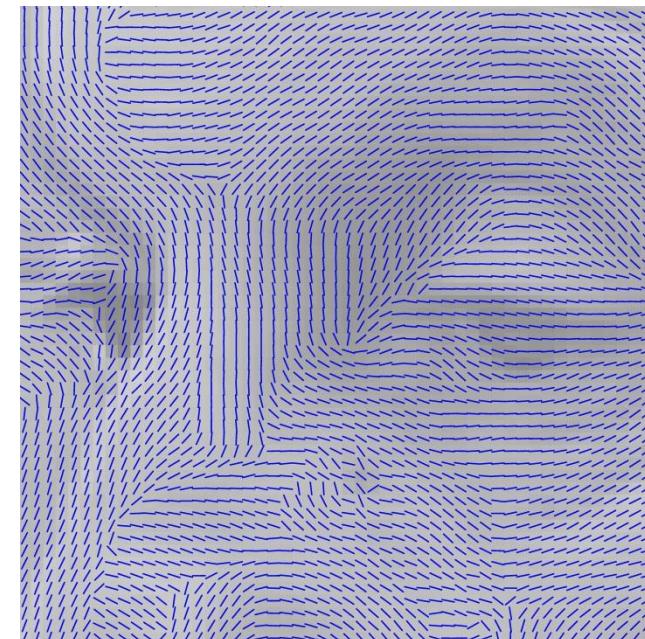


The eigenvectors corresponding to the minor eigenvalues of the structure define a vector field. Typically this field is not smooth:



Smoothed Structure Tensor

Smoothing the structure tensor prior to eigenanalysis with a Gaussian filter removes discontinuities in the vector field:



Smoothed Structure Tensor

Eigenvector field of the smoothed structure tensor is similar to the edge tangent flow, but allows a more efficient implementation:



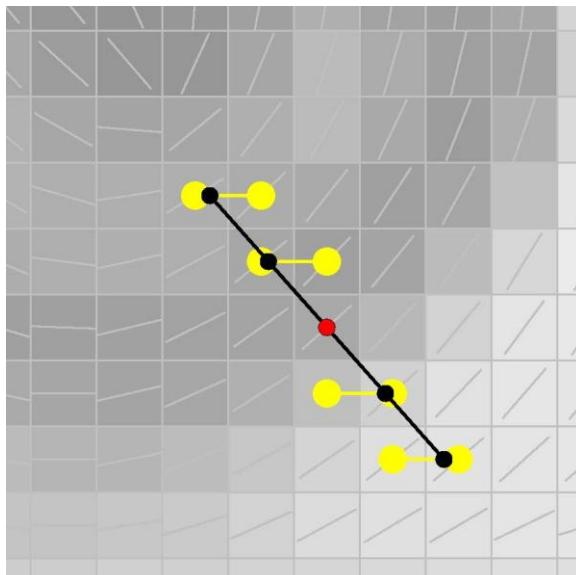
3 iterations of edge tangent flow filter



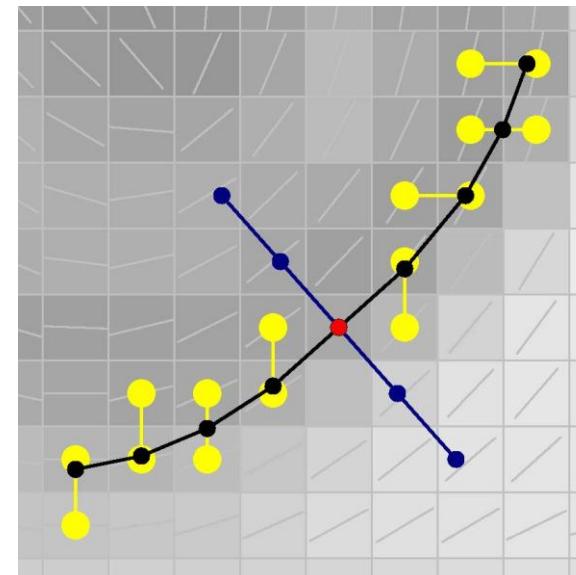
Eigenvector field of the smoothed structure tensor

Split flow-based difference of Gaussians into two passes:

- 1st Pass: one-dimensional DoG in direction of the major eigenvector
- 2nd Pass: smoothing along stream lines defined by minor eigenvector



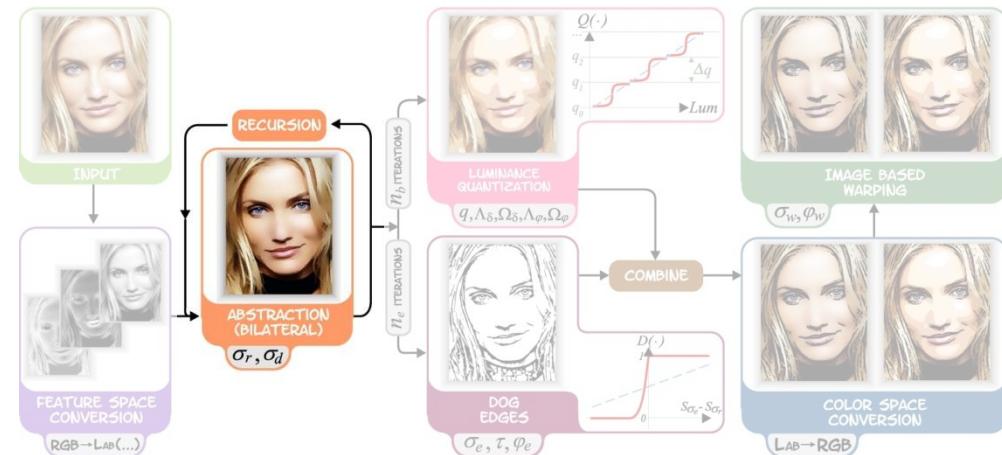
1st Pass



2nd Pass

Bilateral Filter:

- Classical Bilateral Filter
- xy-Separable Bilateral Filter
- Orientation-aligned Bilateral Filter
- Flow-based Bilateral Filter



The bilateral filter is a nonlinear operation that smoothes images while preserving edges:

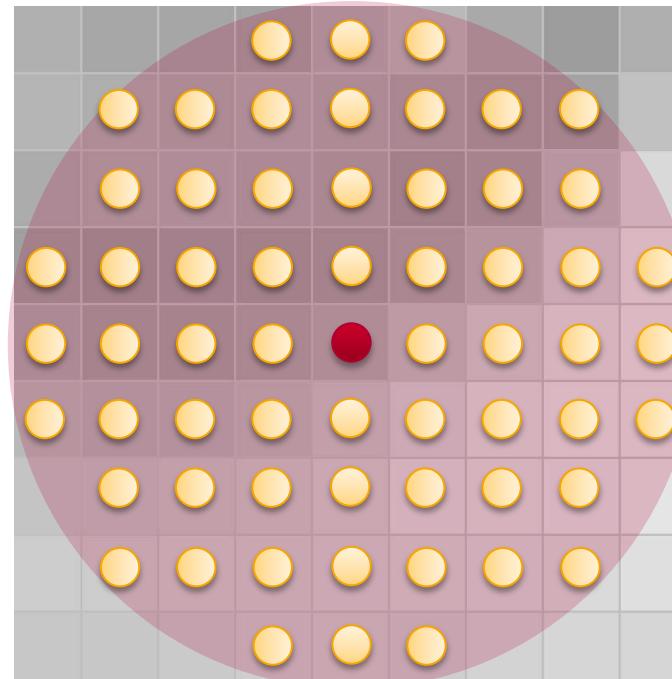
$$\frac{\sum_{x \in B_r(x_0)} f(x) G_{\sigma_d}(|x - x_0|) G_{\sigma_r}(|f(x) - f(x_0)|)}{\sum_{x \in B_r(x_0)} G_{\sigma_d}(|x - x_0|) G_{\sigma_r}(|f(x) - f(x_0)|)}$$

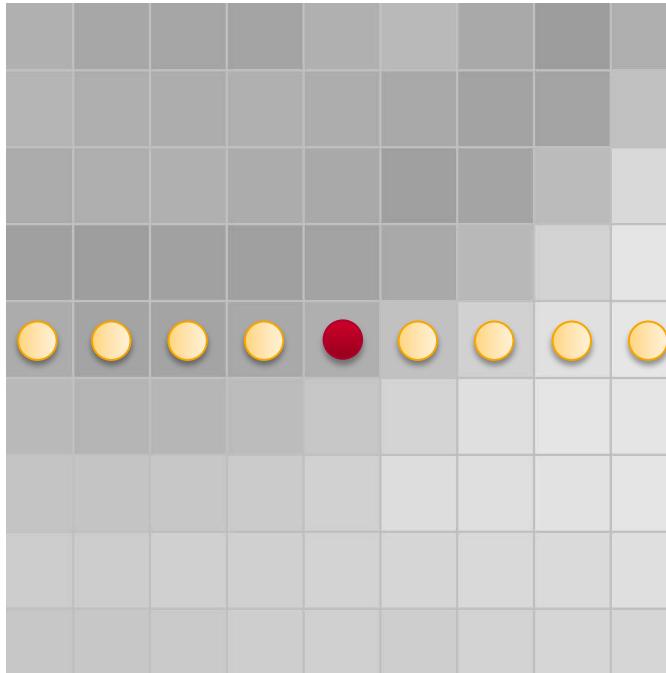
More weight to closer pixel

More weight to pixel with similar color

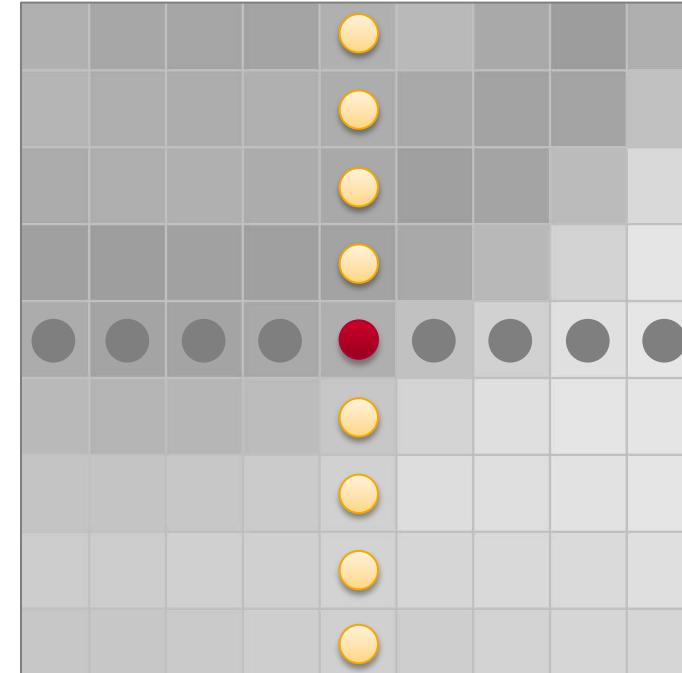
$$G_\sigma(t) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{t^2}{2\sigma^2}\right)$$

The bilateral filter is a powerful tool, but computationally very expensive ($O(r^2)$ per pixel).





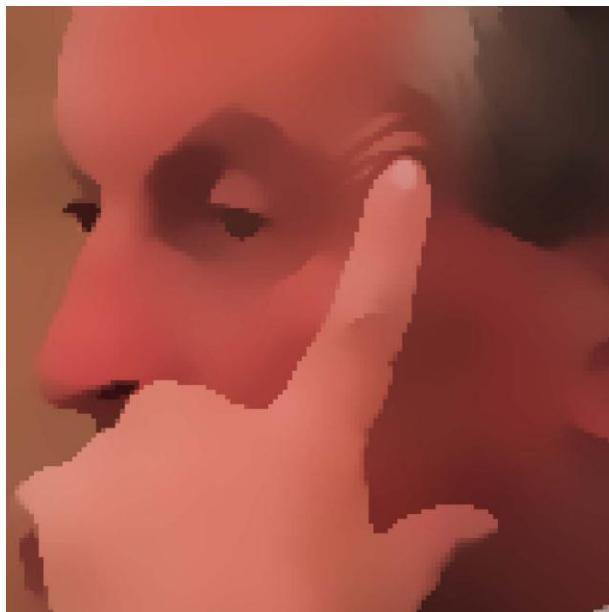
1st Pass



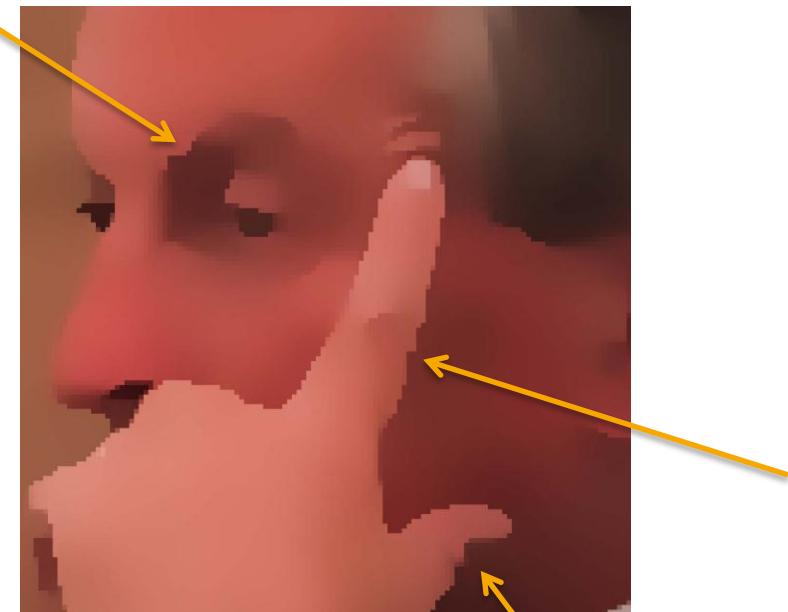
2nd Pass



- Much faster than classical bilateral filter
- But creates noticeable artifacts!

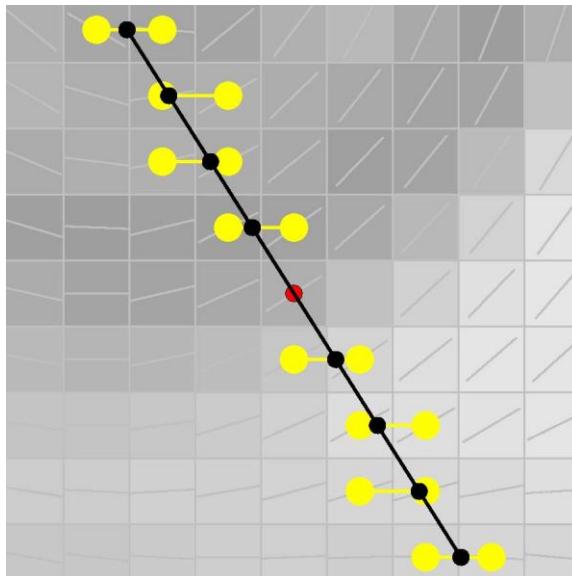


Full kernel bilateral filter

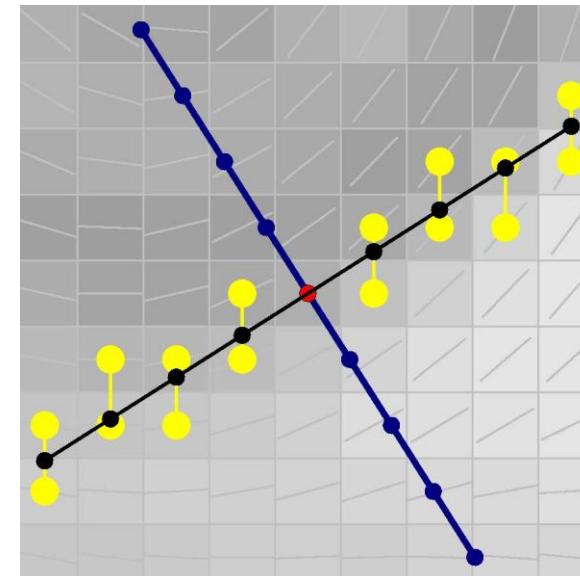


xy-separable bilateral filter

- Align separable bilateral filter to local orientation derived from the smoothed structure tensor



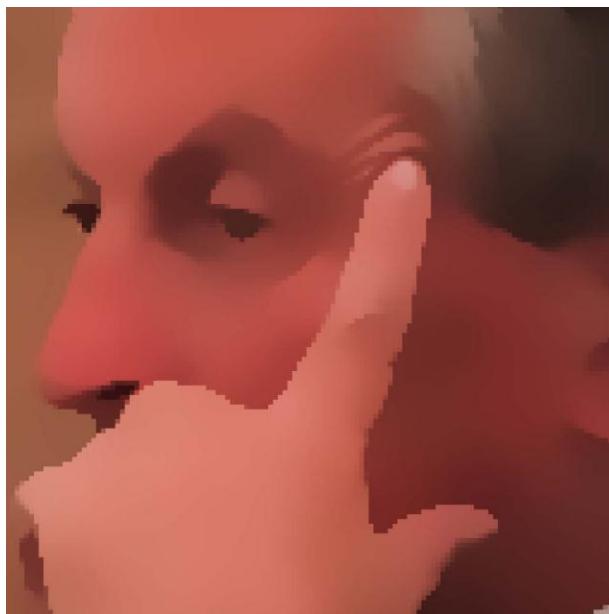
1st Pass



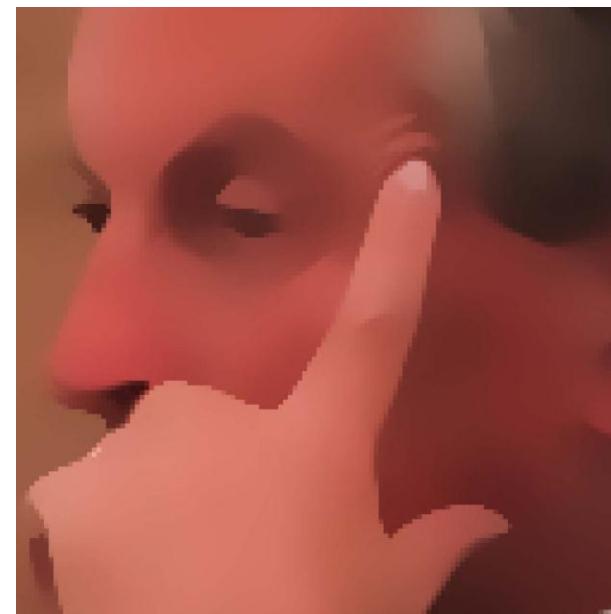
2nd Pass



- Less artifacts. Very well suited for abstraction.



Full kernel bilateral filter

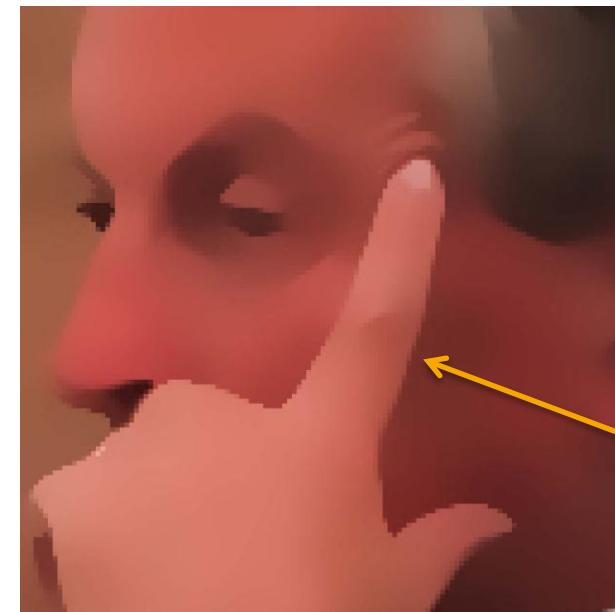
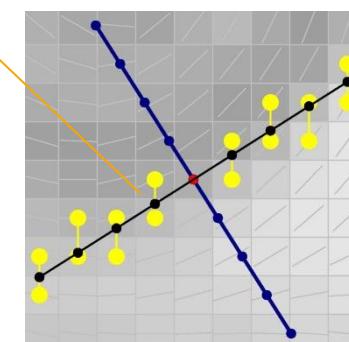
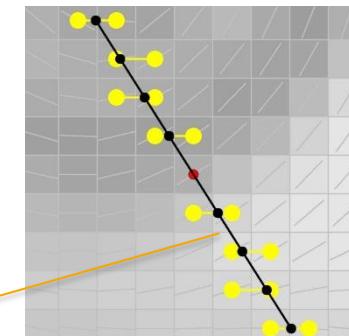


Orientation-aligned bilateral filter

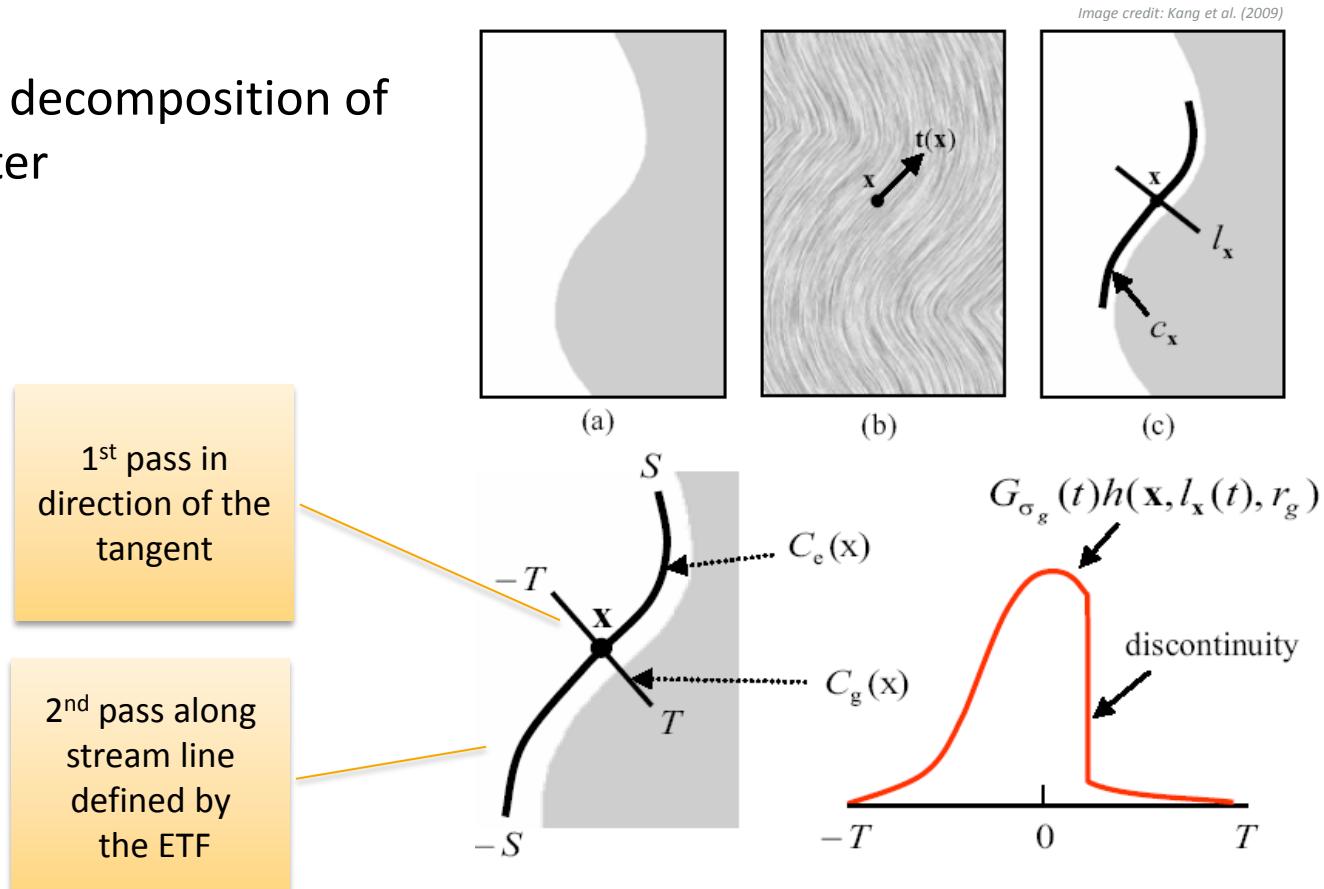


- Linear smoothing of neighboring pixel values creates smooth color boundaries

Unit step size
either along
x- or y-axis



- ETF-guided decomposition of bilateral filter





- Excellent preservation of highly anisotropic image features

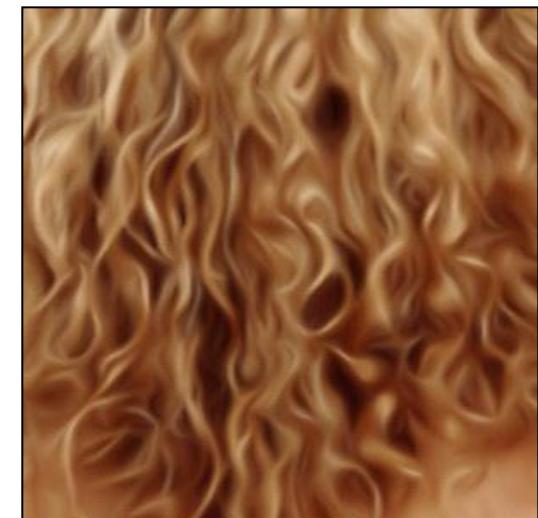
Image credit: Kang et al. (2009)



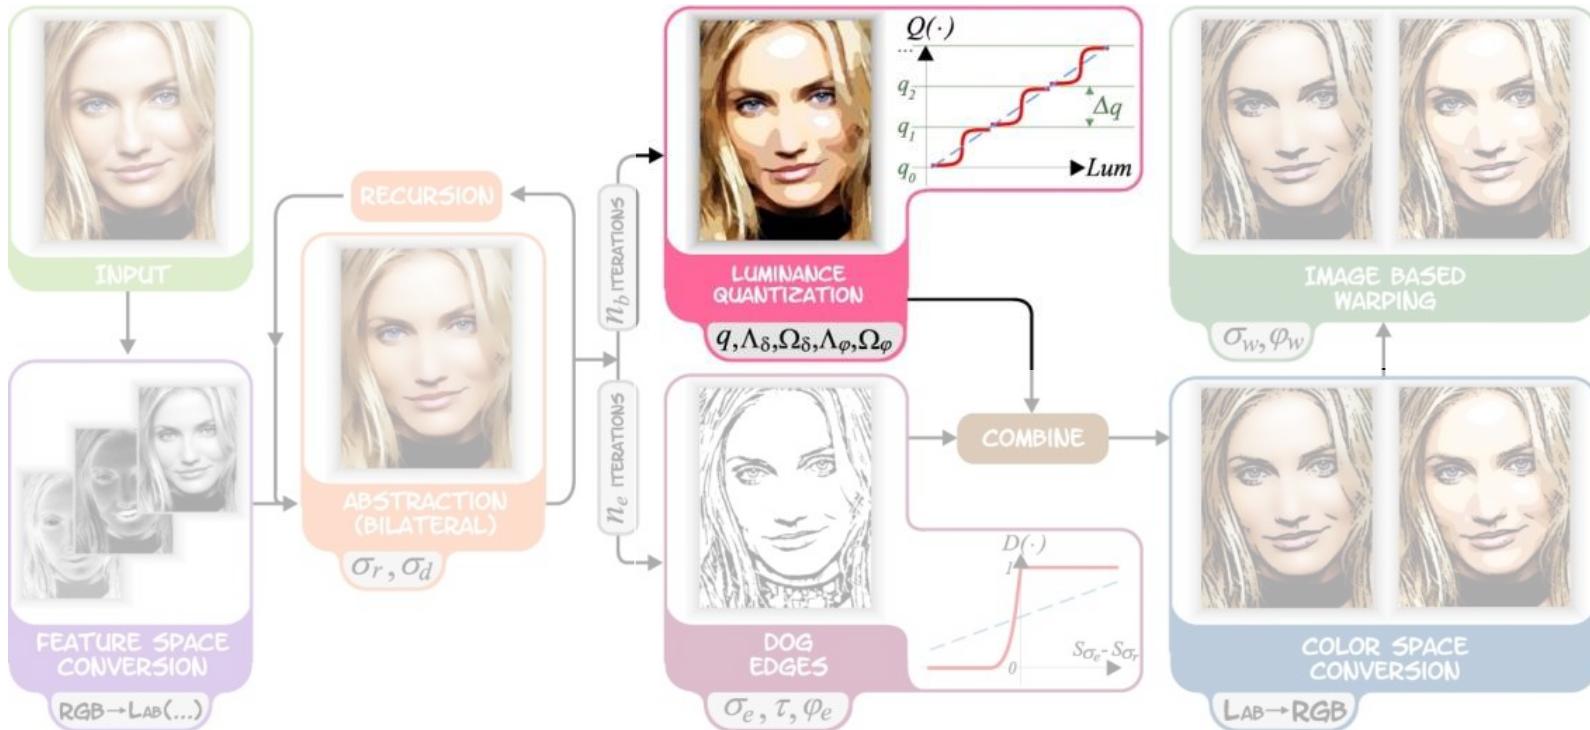
Input image



Bilateral filter



Flow-based bilateral filter





Input



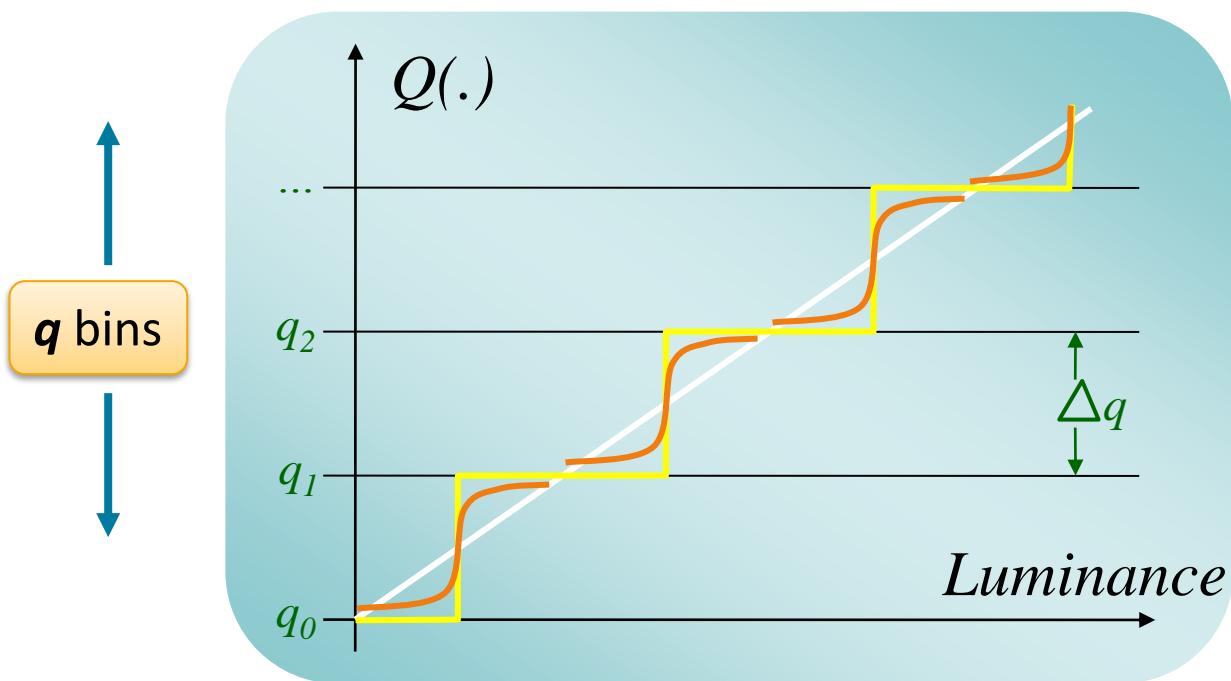
Apply quantization to
luminance channel



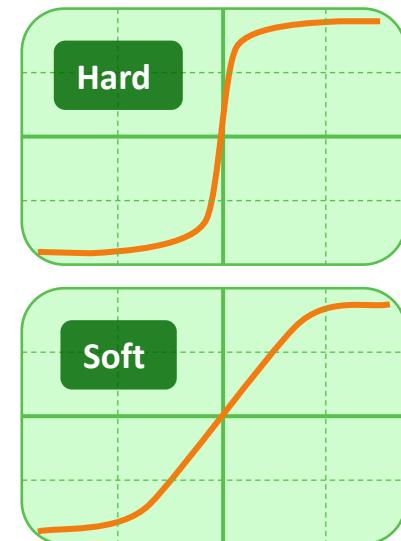
Result

Credit for slide: H. Winnemöller

Luminance Mapping



\tanh
per bin



Credit for slide: H. Winnemöller



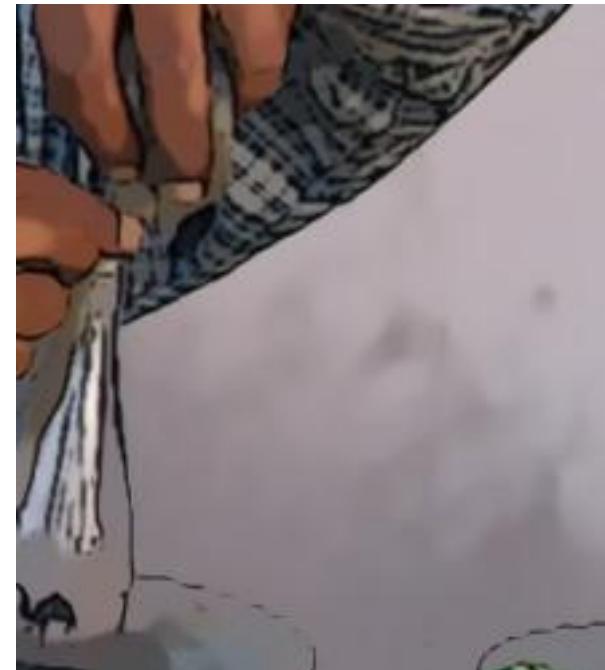
Image credit: Winnemöller et al. (2006)



Abstracted



Sharp Quantization
(Toon-like)

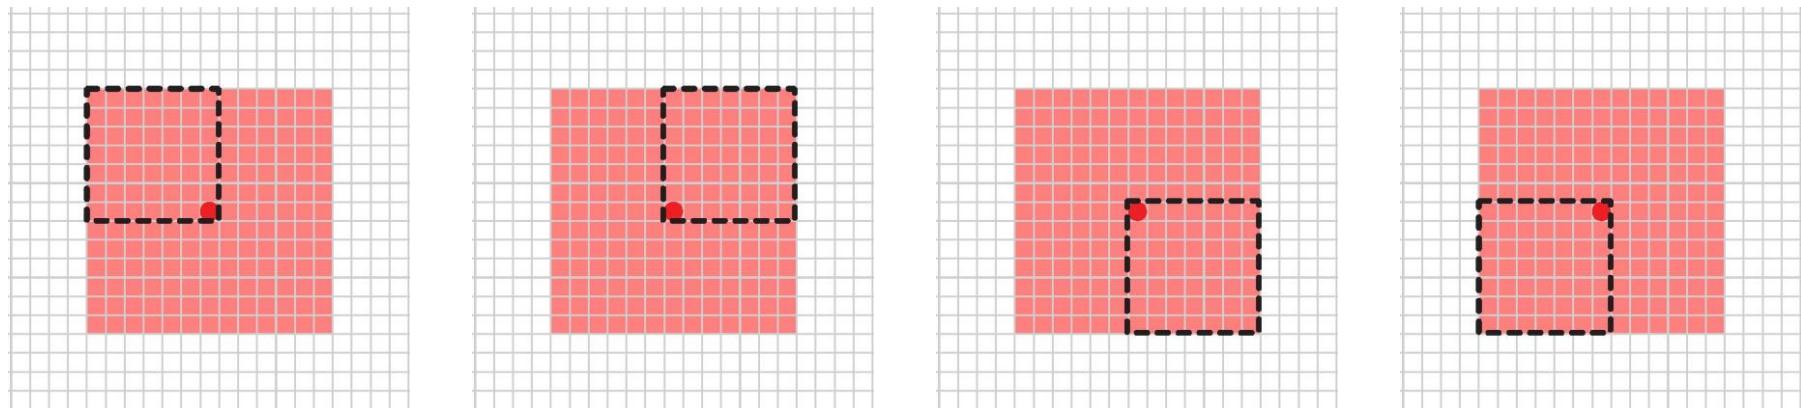


Smooth Quantization
(Paint-like)



Kuwahara Filter:

- Classical Kuwahara Filter
- Kuwahara Filter with Weighting Functions
- Generalized Kuwahara Filter
- Anisotropic Kuwahara Filter
 - Convolution-based Weighting Functions
 - Polynomial Weighting Functions



$$\begin{aligned}W_0 &= [x_0 - r, x_0] \times [y_0, y_0 + r] \\W_1 &= [x_0, x_0 + r] \times [y_0, y_0 + r] \\W_2 &= [x_0, x_0 + r] \times [y_0 - r, y_0] \\W_3 &= [x_0 - r, x_0] \times [y_0 - r, y_0]\end{aligned}$$

For every subregion W_i calculate the mean

$$m_i = \frac{1}{|W_i|} \sum_{(x,y) \in W_i} I(x, y)$$

and the variance:

$$s_i^2 = \frac{1}{|W_i|} \sum_{(x,y) \in W_i} (I(x, y) - m_i)^2$$

The output of the Kuwahara filter is then defined as the mean of a subregion with minimum variance:

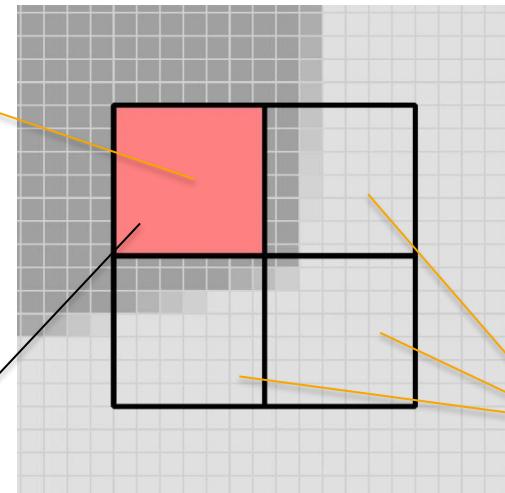
$$F(x_0, y_0) := m_k, \quad k = \operatorname{argmin}_{i=0,\dots,3} s_i$$

Kuwahara filter for a corner

Region with **small variance**. All pixels of this region are similar.

Result is the mean of the region with minimum variance

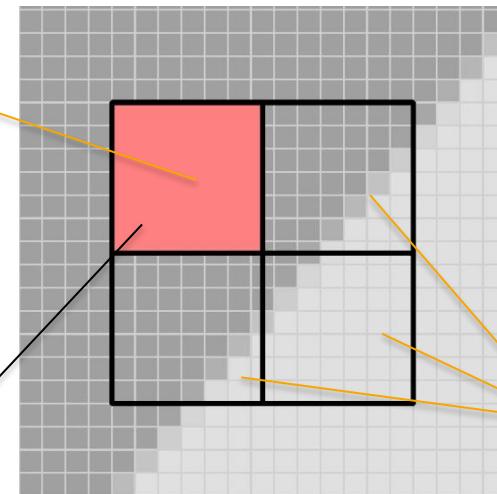
Regions with **high variance**. They all contain pixels from both color regions.



Kuwahara filter for an edge

Region with **small variance**. All pixels of the region lie on the same side of the edge.

Result is the mean of the region with minimum variance

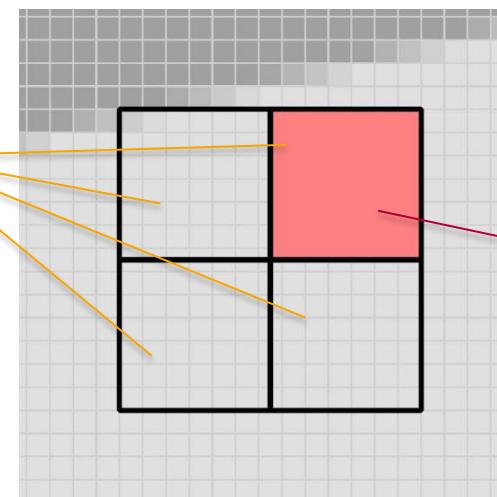


Regions with **high variance**. They all contain pixels from both sides of the edge.



Kuwahara filter for an homogenous neighborhood

All regions have
similar variance!

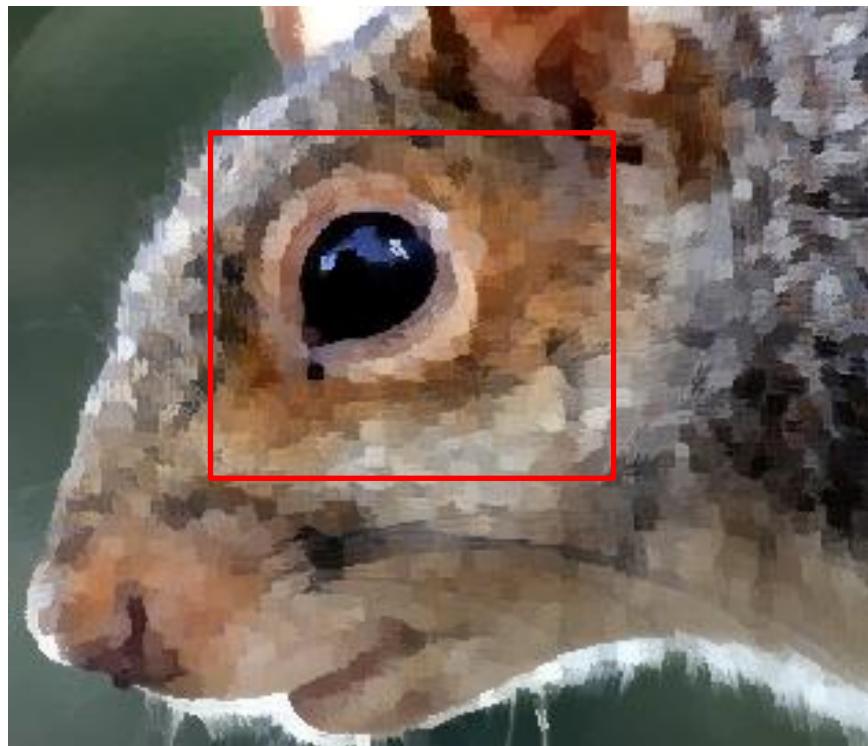


Selection of the
region with
minimum variance
is highly unstable.
For example if
noise is present.

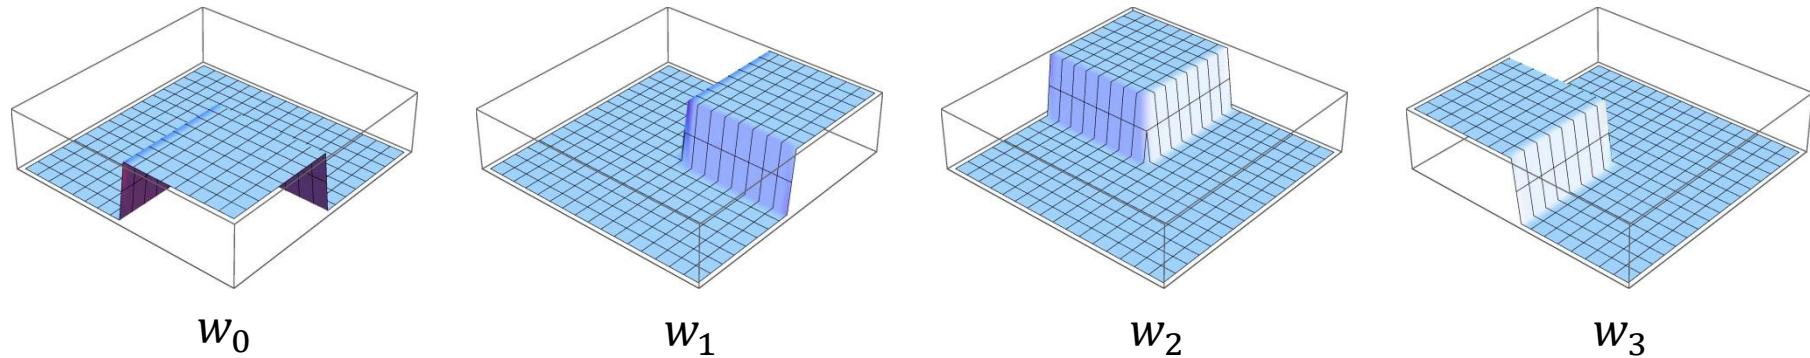


Original image by Keven Law@flickr.com

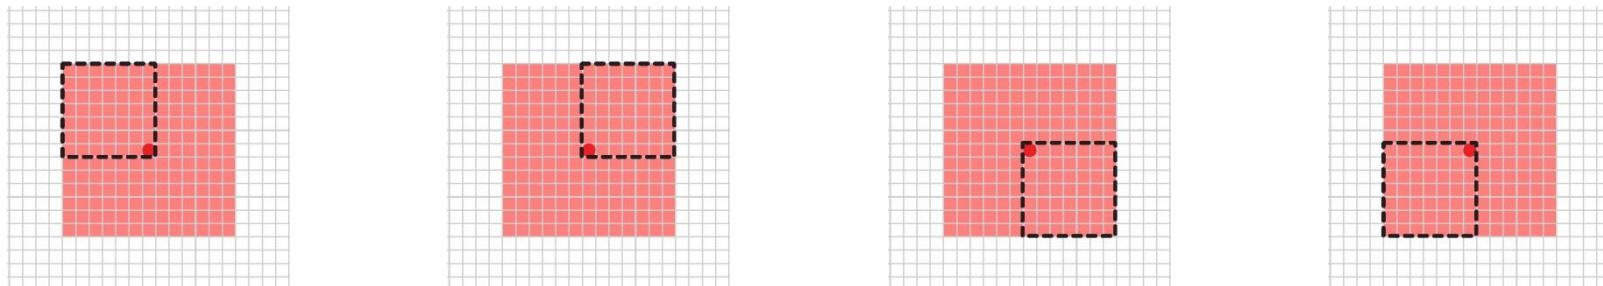




Kuwahara Filter with Weighting Functions



$$w_i(x, y) = \begin{cases} 1 & \text{if } (x, y) \in W_i \\ 0 & \text{otherwise} \end{cases}$$



Kuwahara Filter with Weighting Functions

Then the mean is given by:

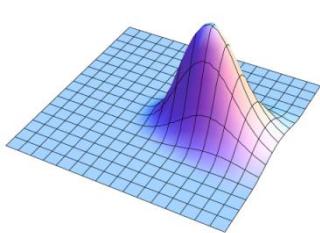
$$\begin{aligned}m_i &= \frac{1}{|W_i|} \sum_{(x,y) \in W_i} I(x, y) \\&= \frac{1}{|w_i|} \sum_{(x,y) \in \mathbb{Z}^2} I(x, y) \cdot w_i(x - x_0, y - y_0)\end{aligned}$$

And the variance is given by:

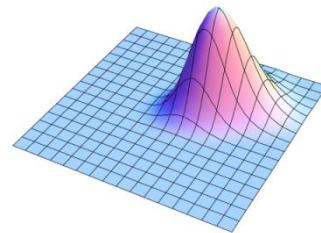
$$\begin{aligned}s_i^2 &= \frac{1}{|W_i|} \sum_{(x,y) \in W_i} (I(x, y) - m_i)^2 \\&= \frac{1}{|w_i|} \sum_{(x,y) \in \mathbb{Z}^2} (I(x, y) - m_i)^2 \cdot w_i(x - x_0, y - y_0)\end{aligned}$$



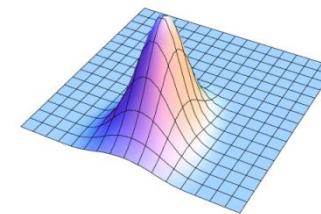
Idea: Create smooth weighting functions over a disc those sum is a Gaussian



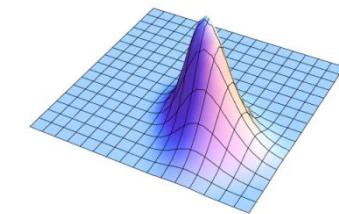
K_0



K_1

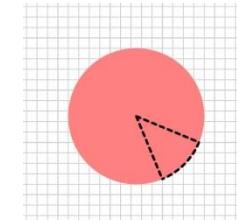
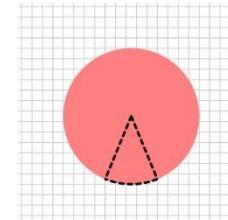
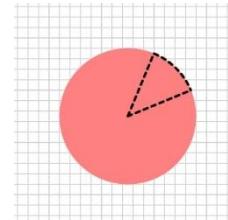
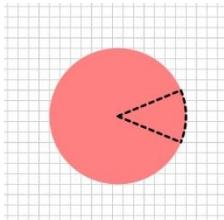


K_6



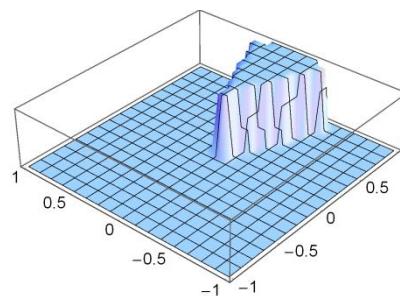
K_7

...

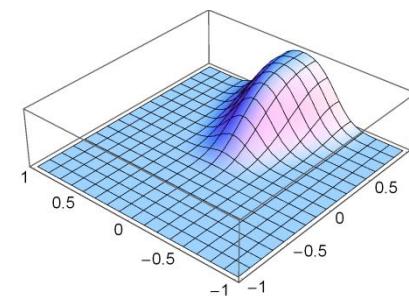


$$\chi_i(x, y) = \begin{cases} 1 & \frac{(2i - 1)\pi}{N} < \arg(x, y) \leq \frac{(2i + 1)\pi}{N} \\ 0 & \text{otherwise} \end{cases} \quad i = 1, \dots, N - 1$$

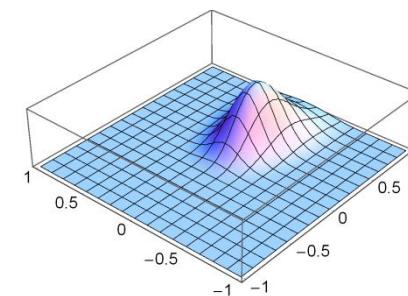
$$\begin{aligned} K_i &= (\chi_i \star G_{\sigma_s}) \cdot G_{\sigma_r} \\ &= K_0 \circ R_{-2\pi i / N} \end{aligned}$$



χ_0



$\chi_0 \star G_{\sigma_s}$



$(\chi_0 \star G_{\sigma_s}) \cdot G_{\sigma_r} =: K_0$

Then the mean is given by:

$$m_i = \frac{1}{|K_i|} \sum_{(x,y) \in \mathbb{Z}^2} I(x, y) \cdot K_i(x - x_0, y - y_0)$$

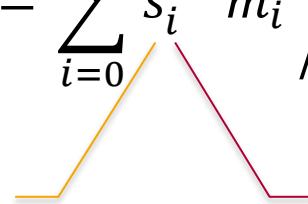
And the variance is given by:

$$s_i^2 = \frac{1}{|K_i|} \sum_{(x,y) \in \mathbb{Z}^2} (I(x, y) - m_i)^2 \cdot K_i(x - x_0, y - y_0)$$

The output of the generalized Kuwahara Filter is now defined by:

$$F(x_0, y_0) = \left/ \sum_{i=0}^{N-1} s_i^{-q} m_i \right/ \sum_{i=0}^{N-1} s_i^{-q}$$

The parameter q is a tuning parameter that controls the sharpness of color boundaries. A typical value is $q = 8$.

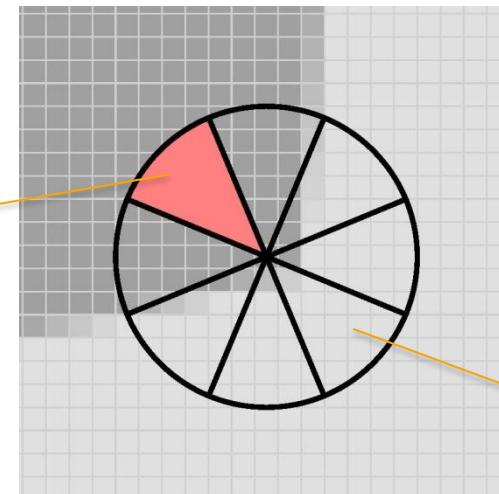
$$F(x_0, y_0) = \frac{\sum_{i=0}^{N-1} s_i^{-q} m_i}{\sum_{i=0}^{N-1} s_i^{-q}}$$


Sectors low high variance:
 $s_i \rightarrow 0 \Rightarrow s_i^{-q} \rightarrow \infty$
(i.e. most influence to sum)

Sectors with high variance:
 $s_i \gg 0 \Rightarrow s_i^{-q} \approx 0$
(i.e. almost no influence to sum)

Generalized Kuwahara filter for a corner

Sector with **small variance**. All pixels of this sector are similar. This sector contributes most to the final result

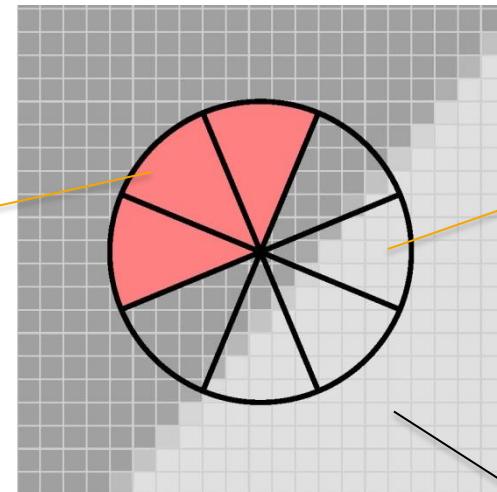


Sectors with **high variance**. They all contain pixels from both color regions. These sectors have almost no influence.

Generalized Kuwahara filter for an edge

Multiple sectors with **small variance**.

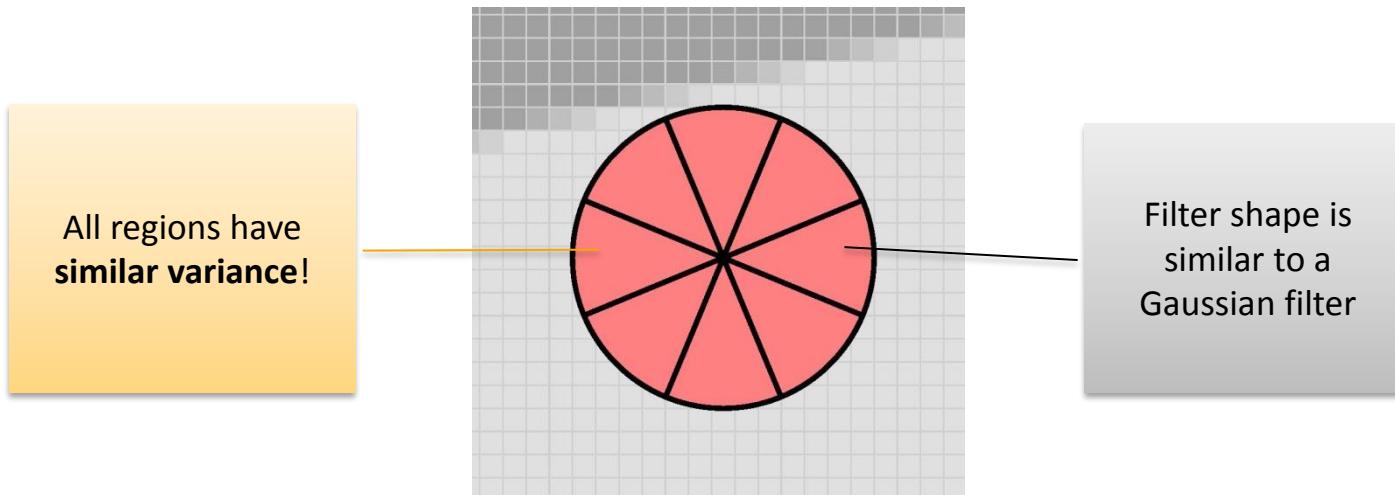
All pixels of the sectors lie on the same side of the edge. Result is a weighted sum of the (weighted) mean values of the sectors.



Regions with **high variance**. They all contain pixels from both sides of the edge. These sector have almost no influence.

Filter shape is similar to a truncated Gaussian

Generalized Kuwahara filter for an homogenous neighborhood





Generalized Kuwahara Filter

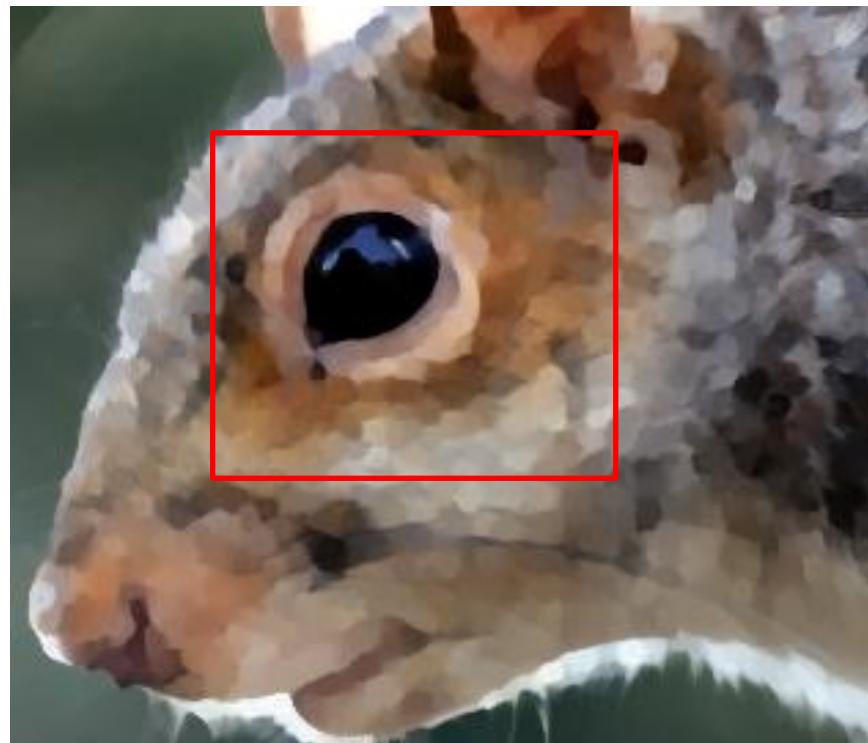


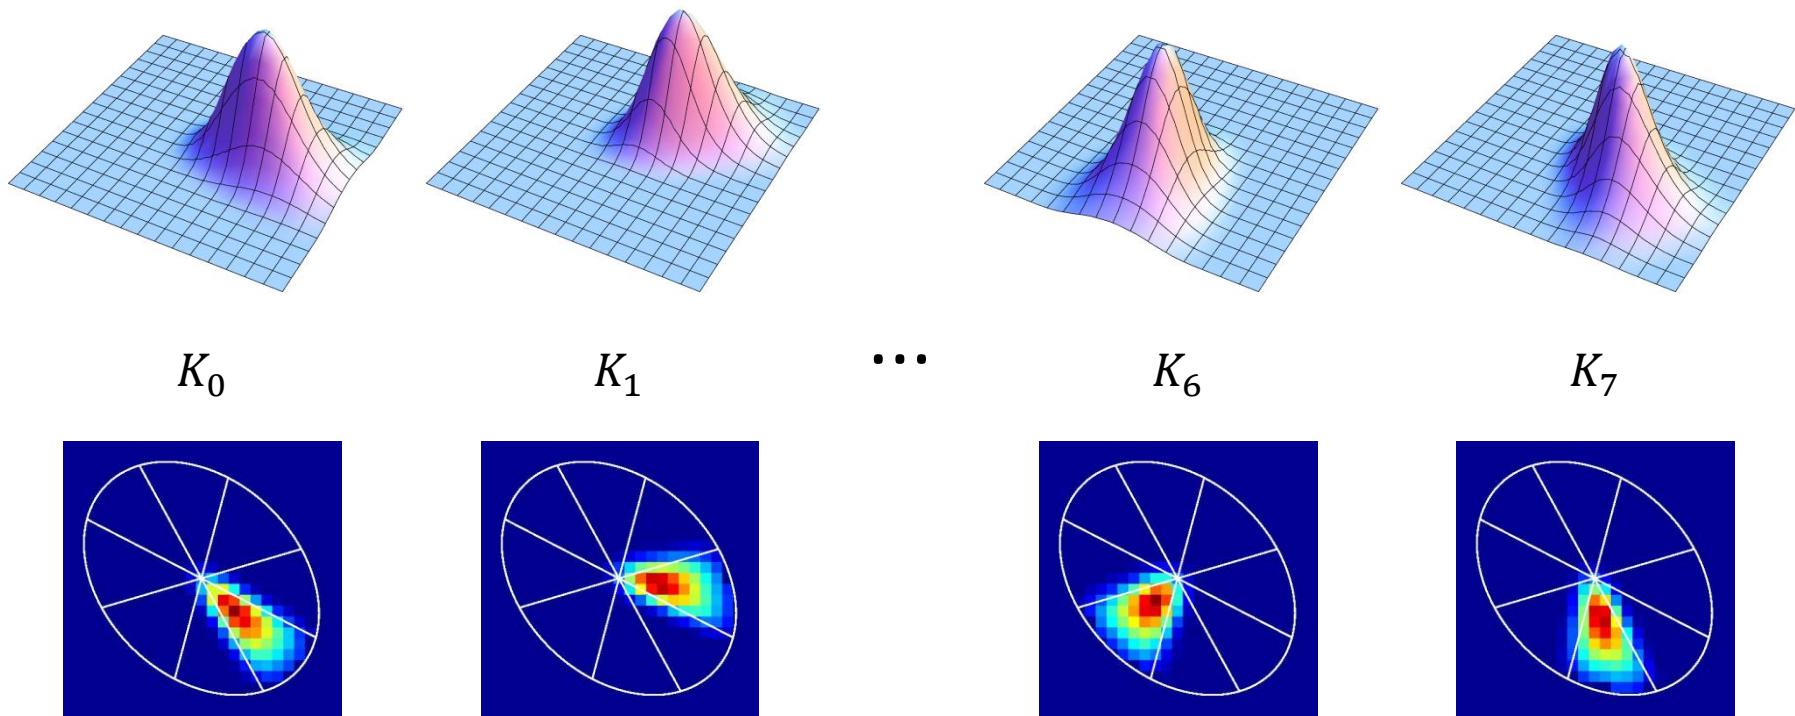
Original image by Keven Law@flickr.com

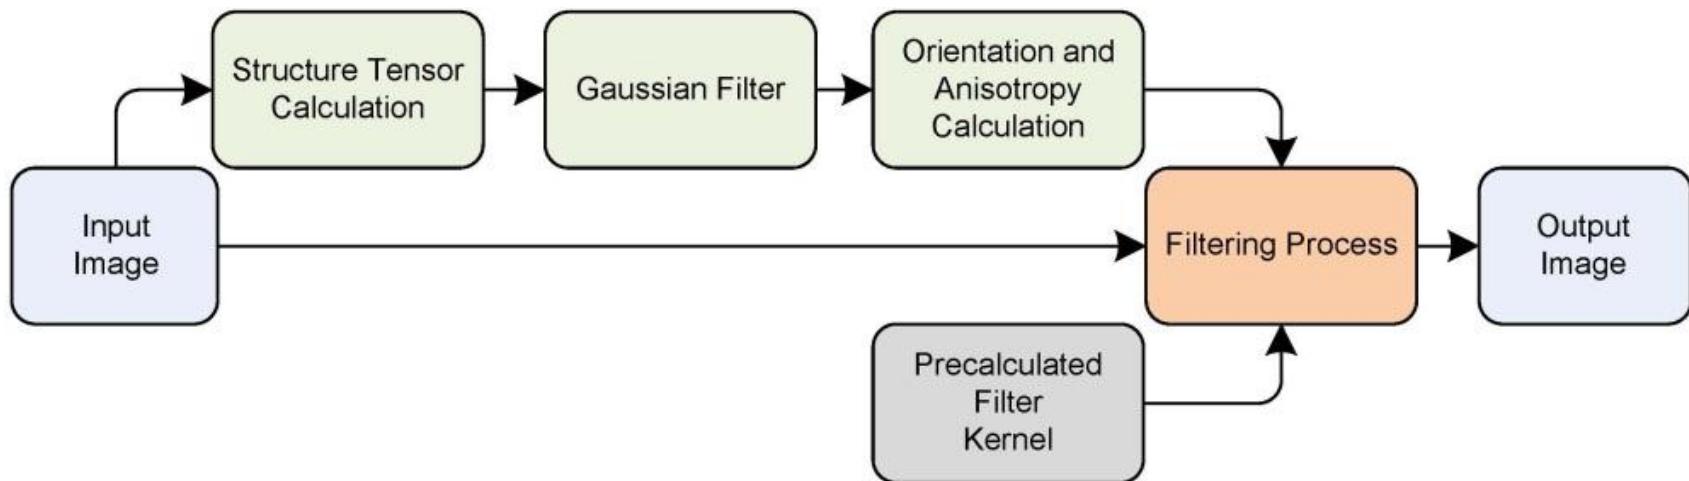




Generalized Kuwahara Filter





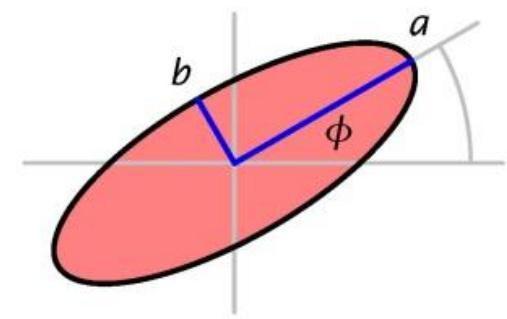


Elliptic filter shape (Pham, 2006)

$$a = \frac{\nu + A}{\nu} \quad b = \frac{\nu}{\nu + A}$$

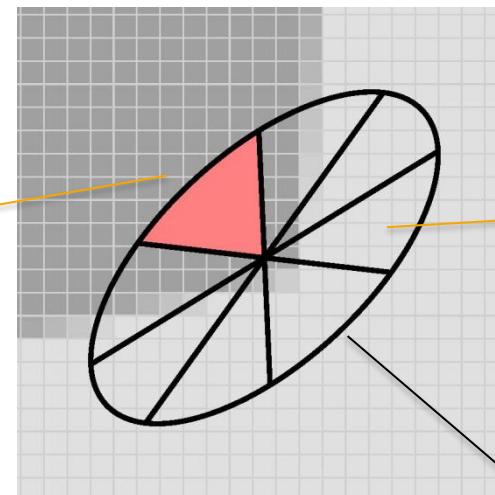
Here, $A \in [0,1]$ denotes the anisotropy measure derived from the structure tensor.

$\nu \in (0, \infty)$ is a user parameter that controls the eccentricity of the ellipse. A typical choice is $\nu = 1$.



Anisotropic Kuwahara filter for a corner

Sector with **small variance**. All pixels of this sector are similar. This sector contributes most to the final result

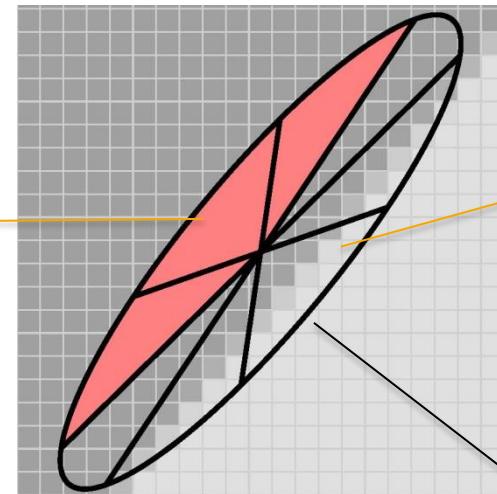


Sectors with **high variance**. They all contain pixels from both color regions. These sectors have almost no influence.

Filter shape is adapted per-pixel

Anisotropic Kuwahara filter for an edge

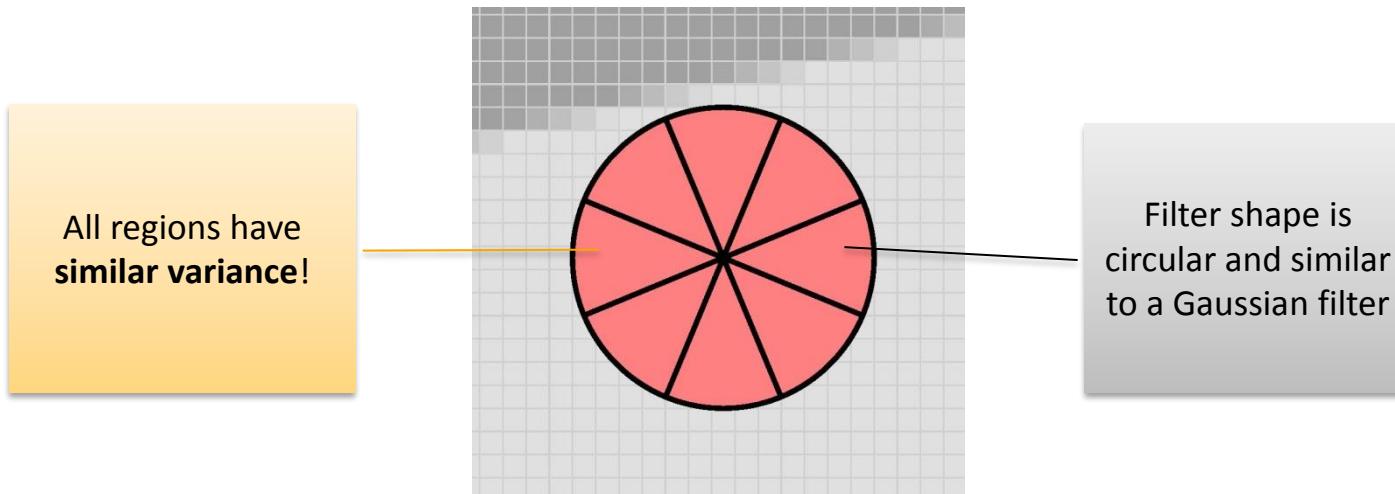
Multiple sectors with **small variance**. All pixels of the sectors lie on the same side of the edge. Result is a weighted sum of the (weighted) mean values of the sectors.



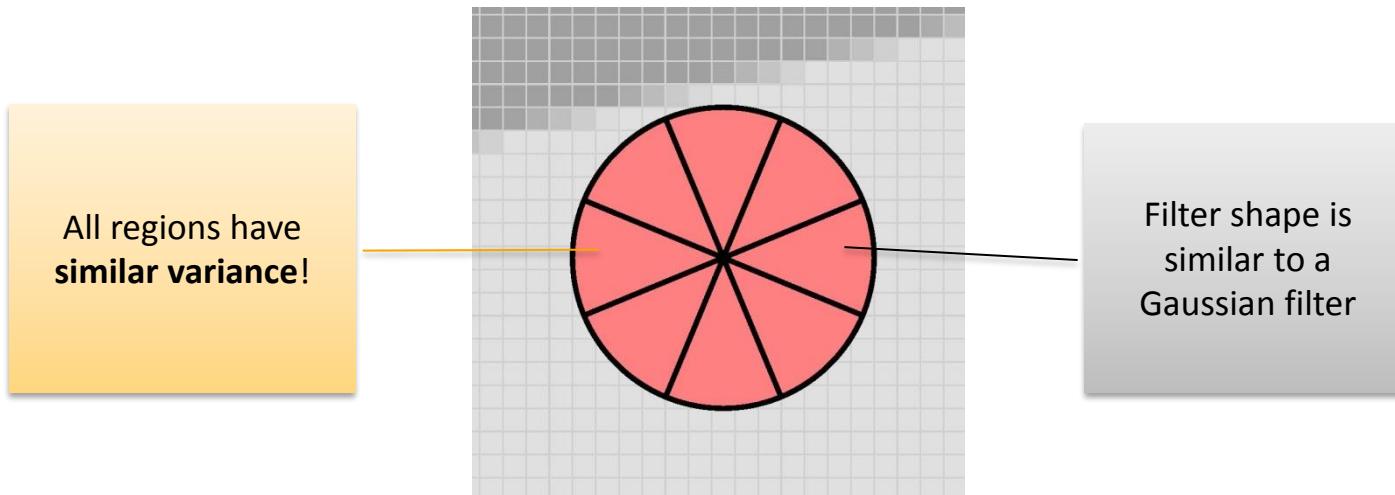
Regions with **high variance**. They all contain pixels from both sides of the edge. These sectors have almost no influence.

Filter shape is adapted to anisotropic image structure

Anisotropic Kuwahara filter for an homogenous neighborhood



Anisotropic Kuwahara filter for a homogenous neighborhood





Anisotropic Kuwahara Filter

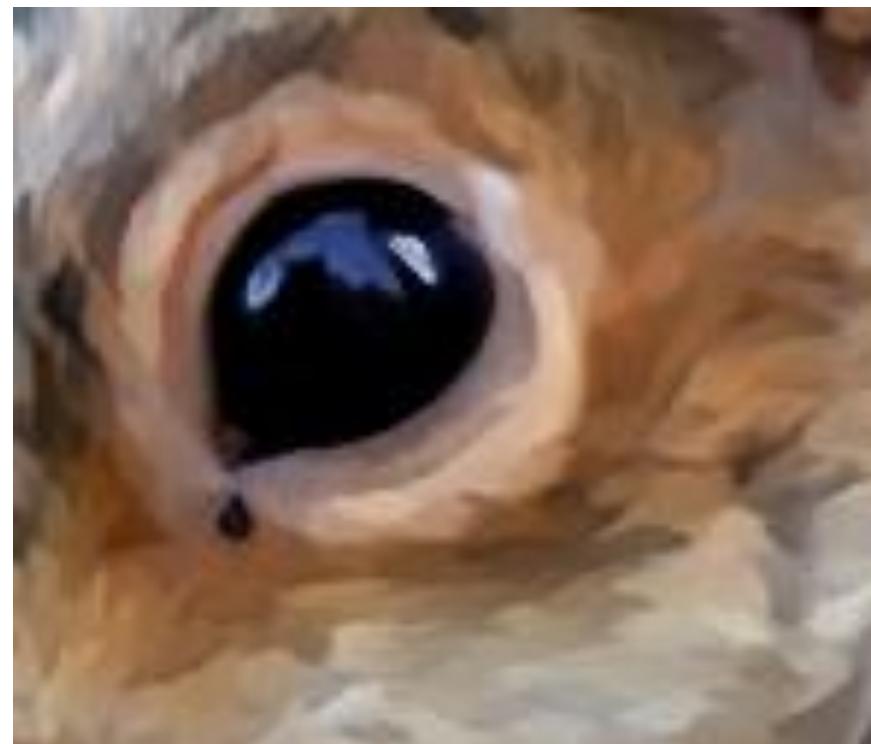
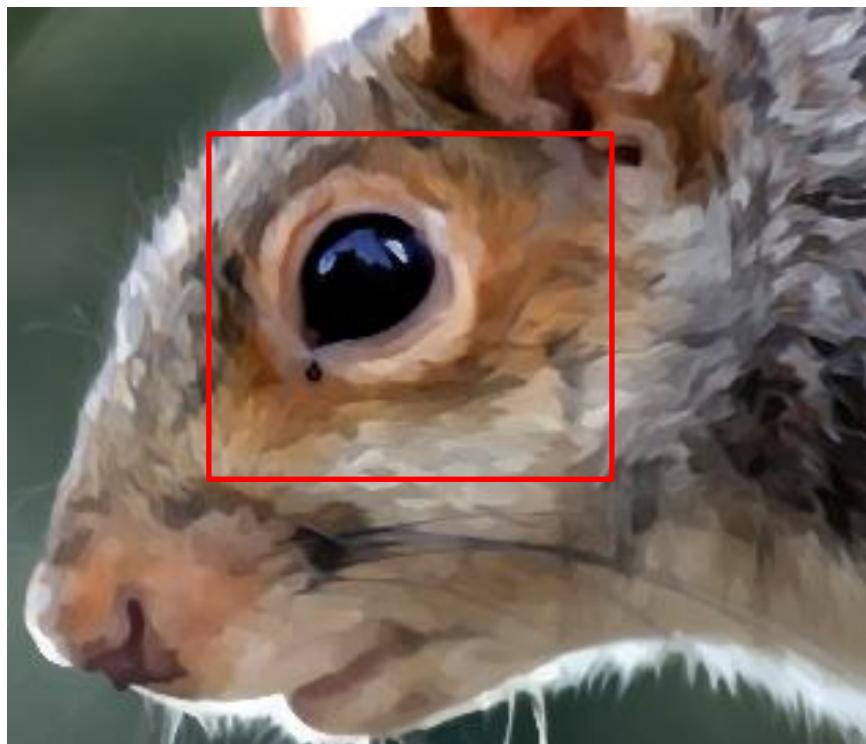


Original image by Keven Law@flickr.com



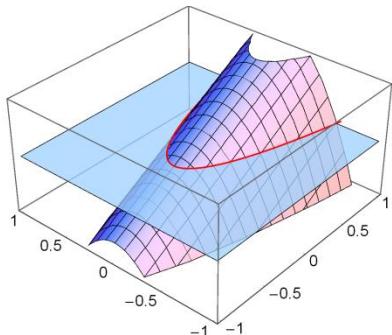


Anisotropic Kuwahara Filter

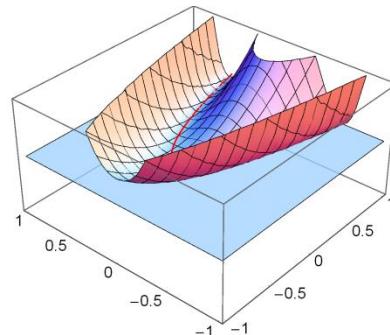


Anisotropic Kuwahara Filter

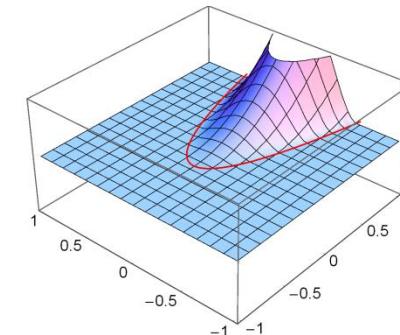
Polynomial Weighting Functions



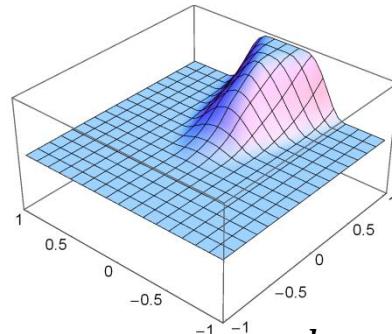
$$(x + \zeta) - \eta y^2$$



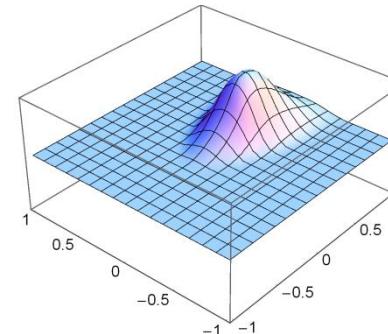
$$((x + \zeta) - \eta y^2)^2$$



$$\max \left\{ 0, ((x + \zeta) - \eta y^2)^2 \right\} := k_0$$

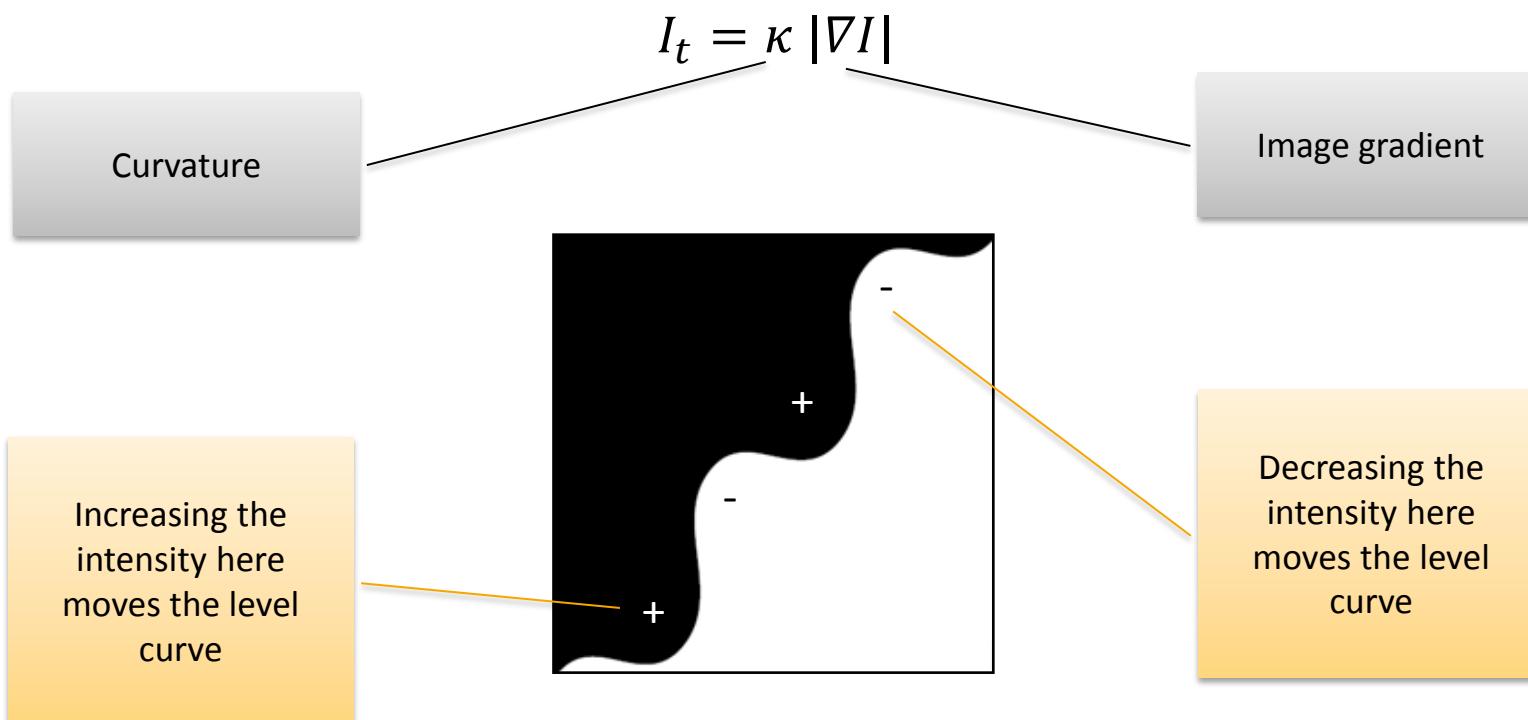


$$\tilde{k}_0(x, y) = \frac{k_0}{\sum_i k_i(x, y)}$$



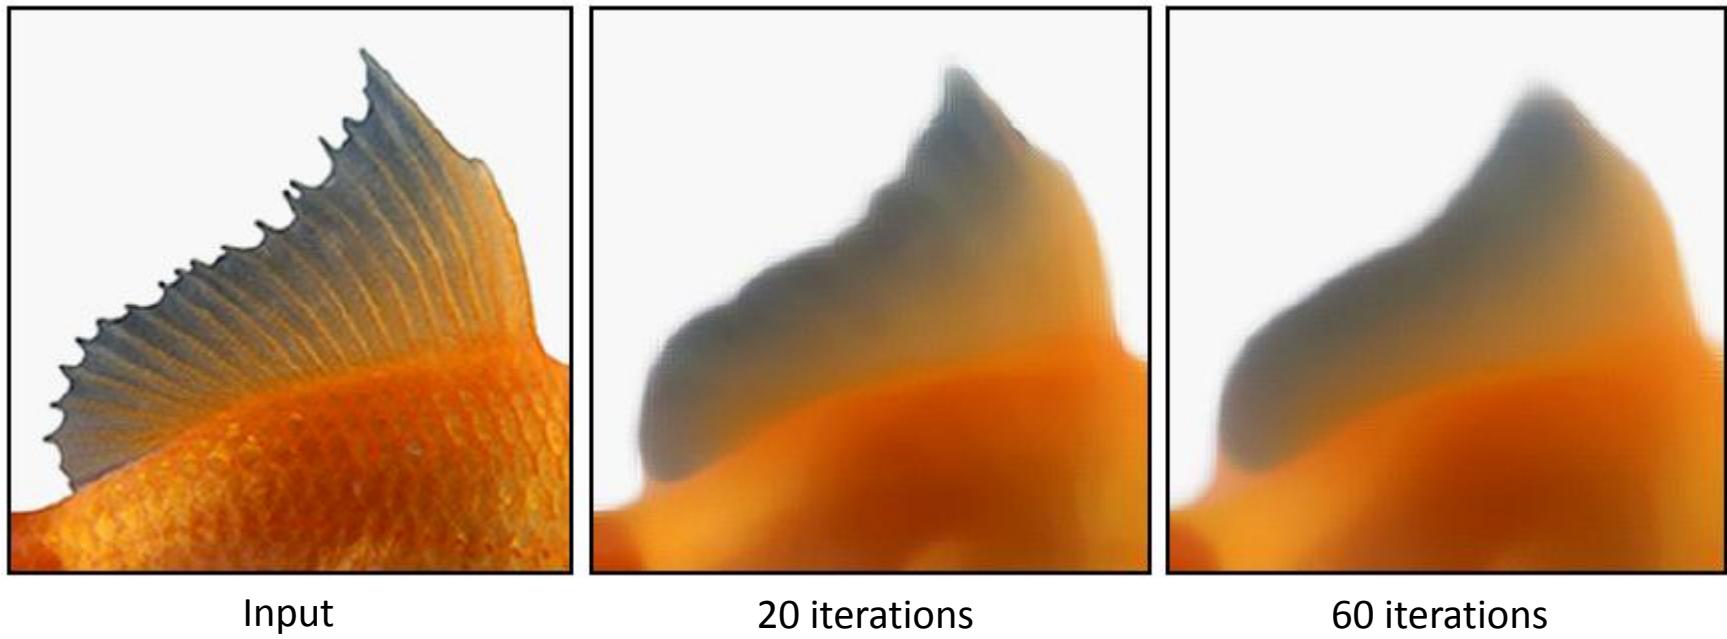
$$\tilde{K}_0 = \tilde{k}_0 \cdot G_{\sigma_s}$$

Mean curvature flow (Alvarez et al., 1992):



Mean curvature flow

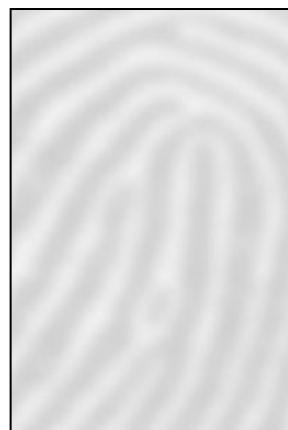
Image credit: Kang & Lee (2008)



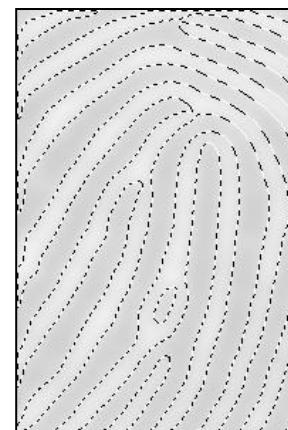
Shock filter (Osher and Rudin, 1990; Alvarez and Mazorra 1994):

In the influence zone of a maximum, the Laplacian ΔI is negative and, therefore, a dilation is performed.

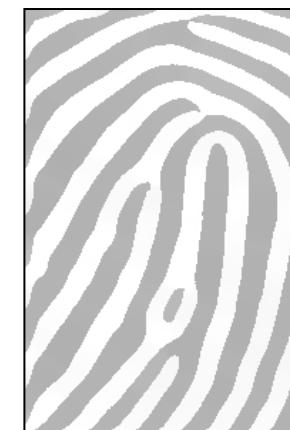
$$I_t = -\text{sign}(\Delta G_\sigma * I) |\nabla I|$$



Input



Influence zones



Output

In the influence zone of a minimum, the Laplacian ΔI is positive, which results in an erosion.

Image credit: Kang & Lee (2008)

Algorithm 1 Image Abstraction by MCF

loop

for 1 to k **do**

$I \leftarrow MeanCurvatureFlow(I)$

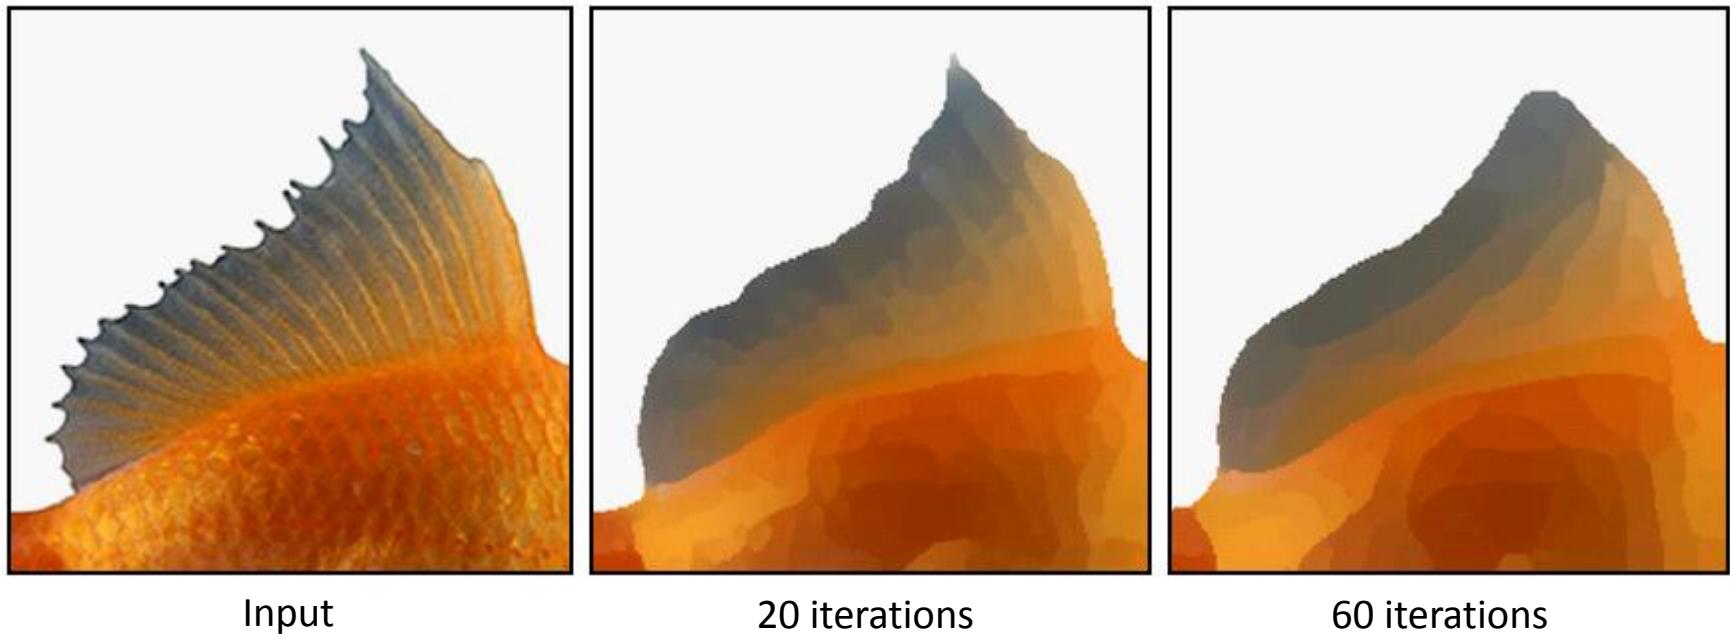
end for

$I \leftarrow ShockFiltering(I)$

end loop

Image abstraction by mean curvature flow

Image credit: Kang & Lee (2008)



Constrained mean curvature flow:

with

$$I_t = s \cdot \kappa |\nabla I|$$

$$s(x) = |t(x) \cdot \nabla I(x)|$$

Stop diffusion if
gradient is not
aligned to edge
tangent flow (ETF)

Edge tangent
flow (ETF)

Image gradient



Algorithm 2 Image Abstraction by CMCF

loop

for 1 to k **do**

$\mathbf{t} \leftarrow TVF(I)$

$I \leftarrow ConstrainedMeanCurvatureFlow(I, \mathbf{t})$

end for

$I \leftarrow ShockFiltering(I)$

end loop

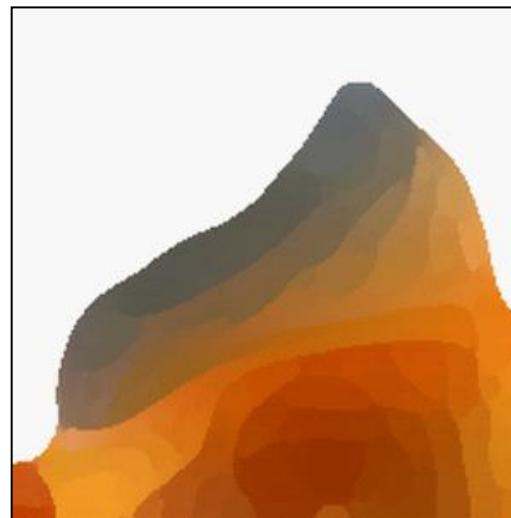


Image abstraction by constrained mean curvature flow

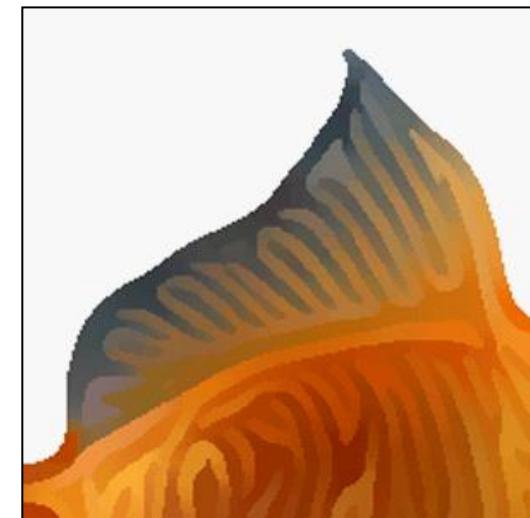
Image credit: Kang & Lee (2008)



Input



60 iterations image
abstraction by MCF



60 iterations image
abstraction by CMCF